

LITERATURE REVIEW

Instructor: Ana Lucia Varbanescu

a.l.varbanescu@uva.nl

Agenda

- A bit of advice on how to approach your literature study
 - Focus on:
 - Goal
 - Method
 - Analysis
 - Writing
- An example


Why literature study?

- To learn the basics of a field
- To grasp the state-of-the-art in a given field
- To understand “gaps” and find new research ideas

In CPP: to extrapolate the knowledge you get in class (fundamentals/basics) to the state-of-the-art analysis

What is the goal ?

We aim to gain (and provide) new knowledge from a set of relevant (and sufficient) publications on a specific topic.



What is the goal?

*We aim to gain (and provide) **new knowledge** from a set of **relevant (and sufficient) publications** on a **specific topic**.*


- Relevant?
 - High quality / timely / ...
- Sufficient?
 - Enough to allow for a conclusion
- Sub-domain?
 - Focused domain of interest
- New knowledge?
 - New information, combining facts *across the papers*

What is the “recipe” ?

We aim to gain (and provide) new knowledge from a set of relevant (and sufficient) publications on a specific topic.





Important steps:

- Select topic
- Select relevant papers
- Read & reflect
- Analyse
- Put it all together = report



SELECT TOPIC

Not.



Bad example:
"8 papers about parallel computing"

Problem?
Domain is too large, papers are not
easily "connected".


Consequence ?
Cannot get sufficient papers for your
review.

Better: domain

Bad-ish example:
"5 papers about GPU computing"

Problem?
Domain is still too large.

Consequence ?
Papers are very diverse. Difficult to
connect and/or compare.

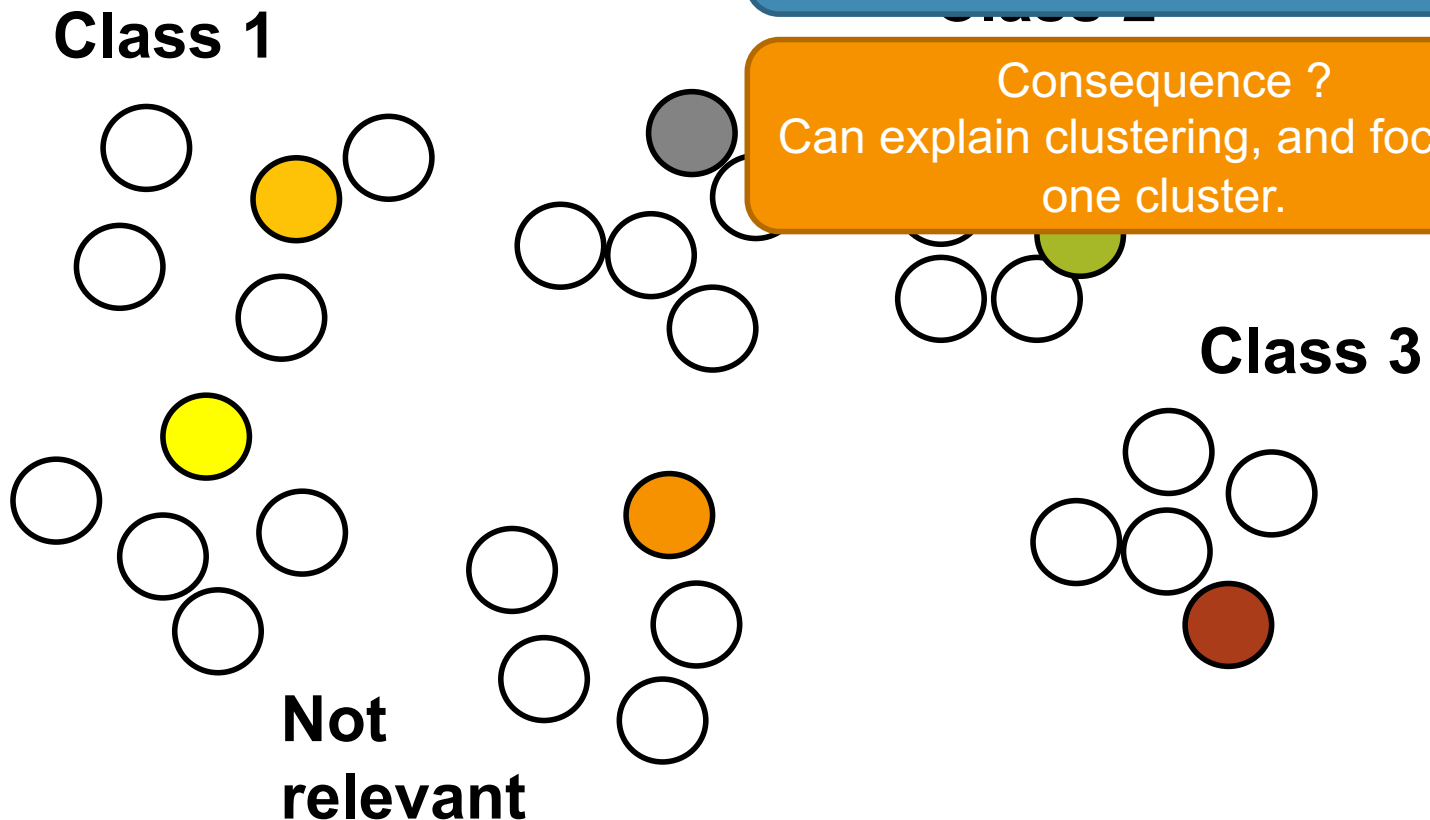


Even Better: cluster, I

Getting better:
"Applications running on GPUs"

Advantage?
Domain is still large, but can be
clustered!

Consequence ?
Can explain clustering, and focus on
one cluster.



Finally, if possible: a summary...

	Papers	Feature 1	Feature 2
This is your domain now!	2	Explain	Explain
Class 2	5	Explain	NA
Class 3	1	Explain	Explain
...	...	--	--

... and a research goal/question


- What are the advances in GPUs for gaming ?
- How has performance changed in ... ?
- What are the common features of GPUs in ... ?
- What are the main features of supercomputers ...

How do you define the domain?

You define the domain: what does your review cover?

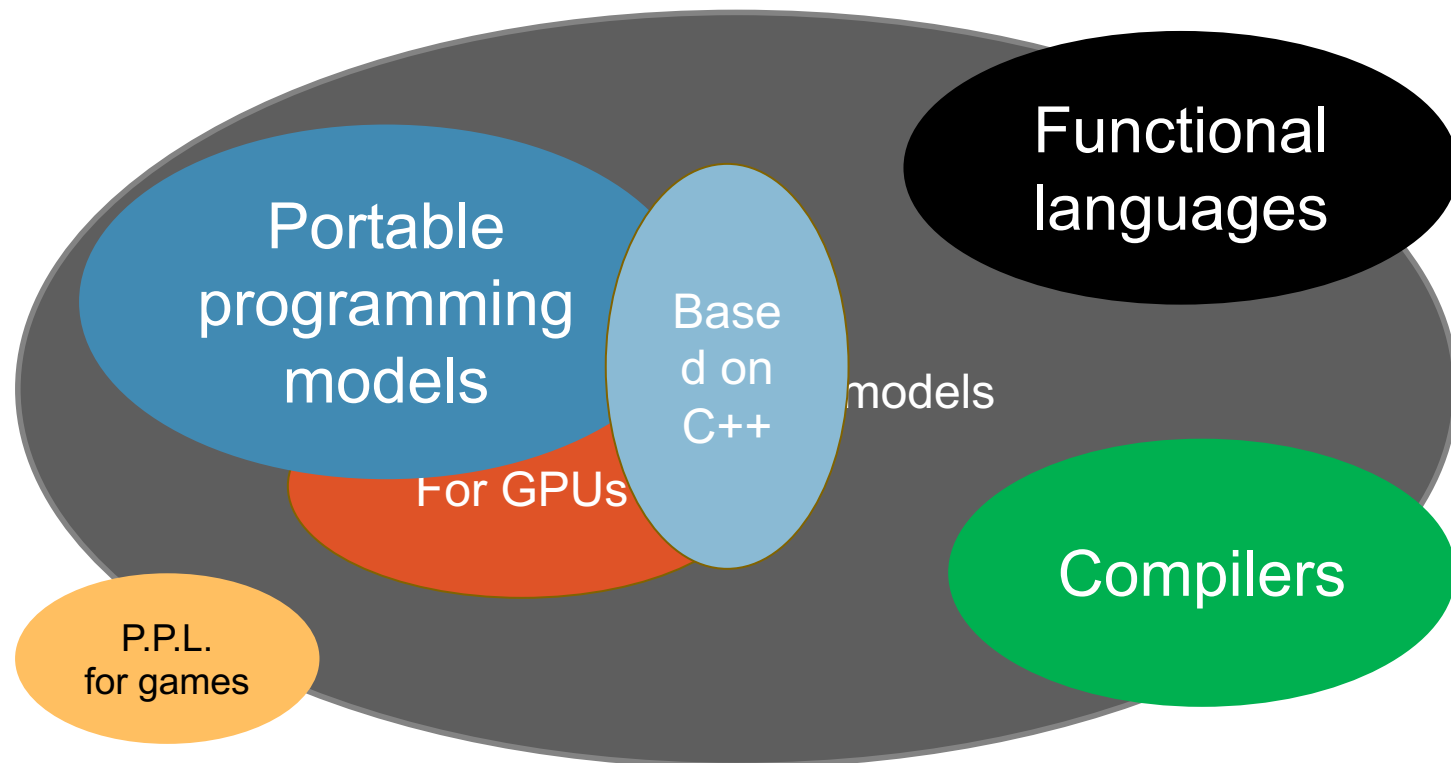
Potential recipe

1. Start from the given domain
2. Select a part that is of interest to you
 - Because you know it
 - Because you can find it in more accessible sites
 - Blogs, Wikipedia, magazines, ...
3. Check for publications
 - Relevant?
 - Sufficient?
4. Repeat ...



Example: Parallel programming models

- A parallel programming model is a programming tool that has specific primitives for parallelism.



Picking a (Sub-)domain

- Try to ensure:
 - Coherency: the domain is easy to define/explain
 - Size: there are sufficient papers
 - Relevance: there are interesting things to say about it
 - Are you curious about it? Go for it !
 - Comparability
 - Can you find a common “thread” ?

Advice(1): Capture your (sub)domain in the research question and the title...

Advice(2): Start early, ask for advice if needed.

Let's practice

- List of topics
 - Football
 - Bicycles
 - Amdahl's law
 - Supercomputers
- Provide a sub-domain and a research question.

Let's try to do this in groups in 5 minutes!

SELECT PUBLICATIONS

Publications

- Multiple sources to search:
 - Google Scholar
 - Unmoderated
 - Metric-based
 - DBLP
 - Moderated
 - ACM Digital Library/IEEE Explore
 - Moderated, but biased
 - Scopus, WoS
 - Year-based searching, chronologies are easier to find
- Please avoid:
 - Blogs, magazines, Wikipedia as main sources
 - Can be used to delimit your domain

Representative publications

- Must touch on the topic you have.
- ... but the results must be ranked:
 - Popularity metrics
 - Paper
 - E.g., Google scholar uses citations
 - Author
 - Objective metrics
 - Age, conference/journal
 - Process metrics
 - Search for X days, read Y papers, ...

... but what about surveys?

- If a survey already exists ...
 - MUST read and analyze
 - Anything missing ?
 - Anything wrong ?
 - Can select some samples of covered topics for more detailed analysis.
 - Can expand the domain.
 - Can propose a different analysis
 - Any dimension the authors did not cover?

READ AND SUMMARIZE

Read papers...

- DFS (=Depth-First Search)
 - Read every paper to completion
 - Analyze each paper individually first
 - Gather similarities and differences.
- BFS (=Breadth-First Search)
 - Read (multiple) papers in layers
 - Analyze the papers together for similarities
 - Analyze differences per paper

Write about what you read

- For every paper you read: 3-5 paragraphs

- Summary
- Positive aspects
- Negative aspects
- Specifics for your analysis
- Relevant related work

} Optional, at times.

- For all papers you read

- Similarities and/or differences
- Build/design/define a taxonomy

Your contribution

- The search process.
- The taxonomy and/or analysis
 - Differences and similarities
 - Insights from different papers.
- The papers' analysis.
- Your general(izing) remarks
 - A sketch of the domain
 - Missing issues
 - Future work ideas

ANALYSE / COMPARE

Analysis

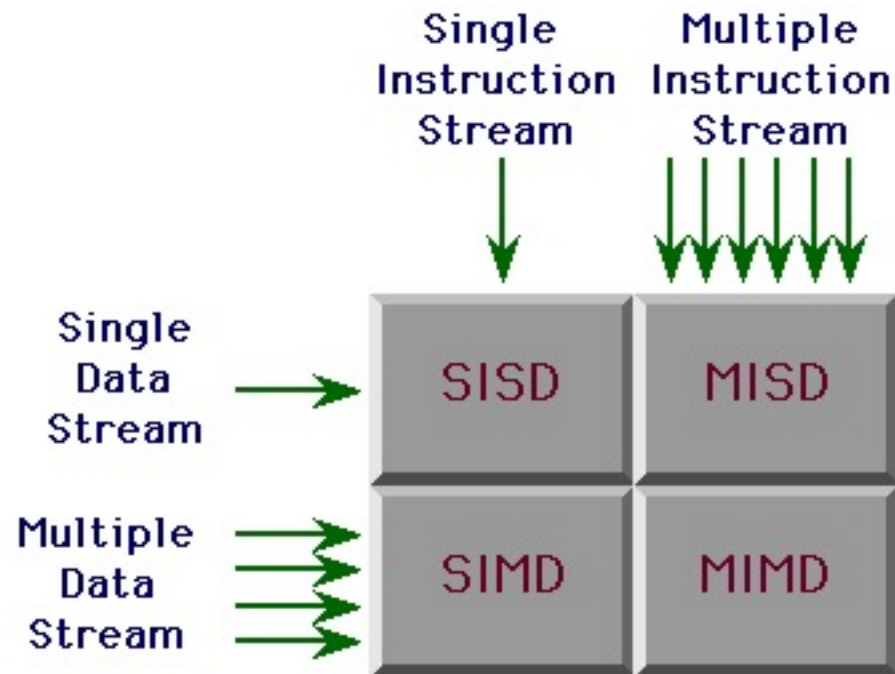
Goal: build a classification/taxonomy of the (sub)domain you have selected.

Taxonomy? Classification?

"Systematic arrangement in groups or categories according to established criteria."

- Classification : non-exhaustive
- Taxonomy : exhaustive
- Main challenge: **determine criteria!**

Example: Flynn's taxonomy



Analysis

Goal: build a classification/taxonomy of the (sub)domain you have selected.

- Bottom-up:
 - Read papers
 - Summarize findings
 - Generalize findings
- Top-down
 - Define a set of expectations/requirements
 - Read papers
 - Discuss coverage of requirements/expectations

Example: bottom-up

*What **are** the common features of GPU algorithms for weather prediction?*

- *Approach:*
 - Read N papers
 - Collect features per paper
 - Compare between papers => defines criteria
- *Advantages:*
 - start from read papers (so, no “pre-processing”)
- *Disadvantages:*
 - Difficult to find common criteria => difficult to compare ...
 - Sparse analysis <= papers have nothing in common ...

Examples: top-down

*What **should be** the main features of GPU algorithms for weather prediction?*

- *Approach:*
 - *Pre-define criteria/features that you WANT to have – “wish-list”*
 - Read N papers
 - Collect features per paper
 - *Compare against the “wish-list”*
- **Advantages:**
 - Always possible to compare
 - Easier to read the papers
- **Disadvantages:**
 - Might find that your wish-list is too optimistic/pesimistic
 - Usually we fix this by refining the wish-list

Example: bottom-up

- GPU energy efficiency models:

	Theoretical/Practical	Use	Power savings	Usable across different architectures
Understanding the Future of Energy Efficiency in Multi-ModuleGPUs	Practical	Measuring	n/a	Yes
GPUWattch: enabling energy optimizations in GPGPUs	Practical	Measuring	14.4% / 11.2%	Yes
GPU Energy Consumption Optimization With a Global-Based Neural Network Method	Practical	Power model	4% - 20%	Yes
An integrated GPU power and performance model	Theoretical	Power model	10.99% / 25.85%	No*
Energy efficiency Analysis of GPUs	Theoretical	Power model	n/a	Yes

Example: top-down

- Top500 architectures

Specifications Supercomp	Rank	Architecture	Cores/ CPU	Clock speed (GHz)	L1 Caches, size in KB (instruction and data)	L2 Cache size	L3 Cache size
Summit	1	Power9	22	3.07	32; 32	512 KB	120 MB
Sierra	2	Power9	22	3.1	32; 32	512 KB	120 MB
Sunway TaihuLight	3	SW26010	260	1.45	32; 32	256 KB	N.A.
Tianhe-2A	4	Ivy Bridge	12	2.2	32; 32	256 KB	30 MB
Frontera	5	Cascade lake	24	2.7	32; 32	1 MB	38.5MB
Piz Daint	6	Haswell	12	2.6	32; 32	256 KB	30 MB
Trinity	7	Haswell	16	2.3	32; 32	1 MB	40 MB
(ABCI)	8	Skylake	20	2.4	32; 32	1 MB	27.5 MB
SuperMUC-NG	9	Cascade lake	24	3.1	32; 32	1 MB	1.375 MB
Lassen	10	Power9	22	3.1	32; 32	512 KB	120 MB

PUT IT ALL
TOGETHER

Step-by-step process

1. Define your (initial) domain (= title)
2. Define approach: bottom-up or top-down
 1. Top-down: define features => sketch taxonomy
3. Refine domain
4. Define search criteria and process
5. Select and read papers
6. Attempt taxonomy/Test taxonomy
 1. Works? Done 😊
 2. Doesn't work -> back to 4, 3, 2.
7. Conclude
 1. Landscape of the (sub-)domain.
 2. Generalized remarks based on taxonomy.
 3. Recommendation for improvement/future work ideas/very interesting points you have found, ...

Table of contents [v1]

- Introduction
 - Motivation
 - Definition of the domain / goals
 - Approach: search process, potential search criteria, ...
- Representative publications
 - Multiple sections, depending emerging classes.
- A taxonomy/classification/sketch
 - Includes generalizing remarks
 - Could be a table to show differences
- Conclusion/Summary & Future Work
 - Findings
 - Reflection on coverage and better approaches

Table of contents [v1]

- Introduction

- Motivation
- Definition of the domain / goals

Approach: search process, potential search criteria

... that's not to say these are the real section titles.

Adapt them to your case / topic / ...

- Includes generalizing remarks
- Could be a table to show differences
- Conclusion/Summary & Future Work
 - Findings
 - Reflection on coverage and better approaches

Table of contents [v2]

- Introduction
 - Motivation
 - Approach: search process, potential search criteria, ...
- Ideal requirements
- Representative publications
 - Possibly multiple sections, depending on how you classify them.
- Coverage of requirements
 - How well does the literature cover the requirements
 - Any paper has them all – highlight
 - Any missing requirement overall – highlight
 - Any observable clustering – highlight
- Conclusion/Summary & Future Work
 - Findings
 - Reflection on coverage and better approaches

Table of contents [v2]

- Introduction
 - Motivation
 - Approach: search process, potential search criteria, ...

- Ideal requirements

... that's not to say these are the real section titles.

Adapt them to your case / topic / ...

- How well does the literature cover the requirements
 - Any paper has them all – highlight
 - Any missing requirement overall – highlight
 - Any observable clustering – highlight
- Conclusion/Summary & Future Work
 - Findings
 - Reflection on coverage and better approaches

AN EXAMPLE

SUMMARY

In summary ...

- Focus on WHAT you are looking for first.
- Decide on how you will approach things
 - Top-down: you know the field
 - Bottom-up: new field to you
- Define and bound your search
- Put your findings in a taxonomy (table, tree, image, classification ...)
- Discuss your findings and highlight things YOU have found
- Summarize/generalize and discuss YOUR ideas for future work
 - In the field or in the analysis. Or both.