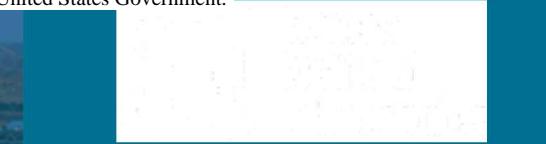
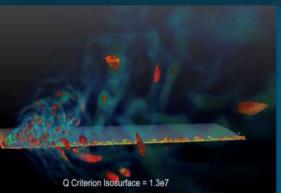


# The Structural Simulation Toolkit and GPGPU-Sim



*PRES E N T E D BY*

Clay Hughes, Sandia National Laboratories

SST Development Team, Sandia National Laboratories

Architecture Accelerator Lab, Purdue University

11/03/2019

# Welcome!

Part 1: Introduction to SST

8:45 - 10:15

SST Overview

*Demos: Running a simple simulation; Enabling statistics; Running in parallel*

SST Element Libraries: A tour

Break

10:15 - 10:45

SST Element Libraries: A tour

*Demos: L2 Cache; Backing store; Adding processors; Adding Memory; Node modeling*

Lunch

12:00 - 1:30

Part 2: GPGPU-Sim

1:30 – 2:30

GPGPU-Sim Overview & New Features

**Part 3: The SST/GPGPU-Sim Integration**

**2:30 - 3:00**

Break

3:00 - 3:30

GPGPU-Sim Exercises

3:30 - 5:00

# SST GPGPU Modeling Capability

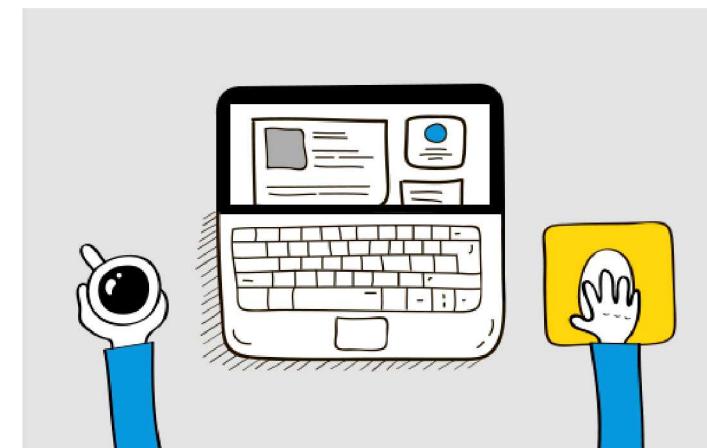


Long history of GPGPU simulators, mostly from academia

- GPGPUSim [Bakhoda, 2009]
- Gem5+GPU [Power, 2014]
- MacSim/Ocelot [Hyesoon, 2012], [Kerr, 2009]
- Multi2Sim v5.0 [Gong, 2017]
- NVArchSim

After a literature review and phone calls with the developers of GPGPUSim and Multi2Sim, the SST team chose GPGPUSim as the integration target

- More amendable to SST integration
  - Software designed around well-defined objects
- More tractable for SST workflows



# GPGPUSim Integration



## GPGPUSim

- Runs *host* code directly
- Runs *device* code on a functional simulator
- Interfaces with programs through a custom library, replacing native CUDA calls
- Simulation split into functional and timing



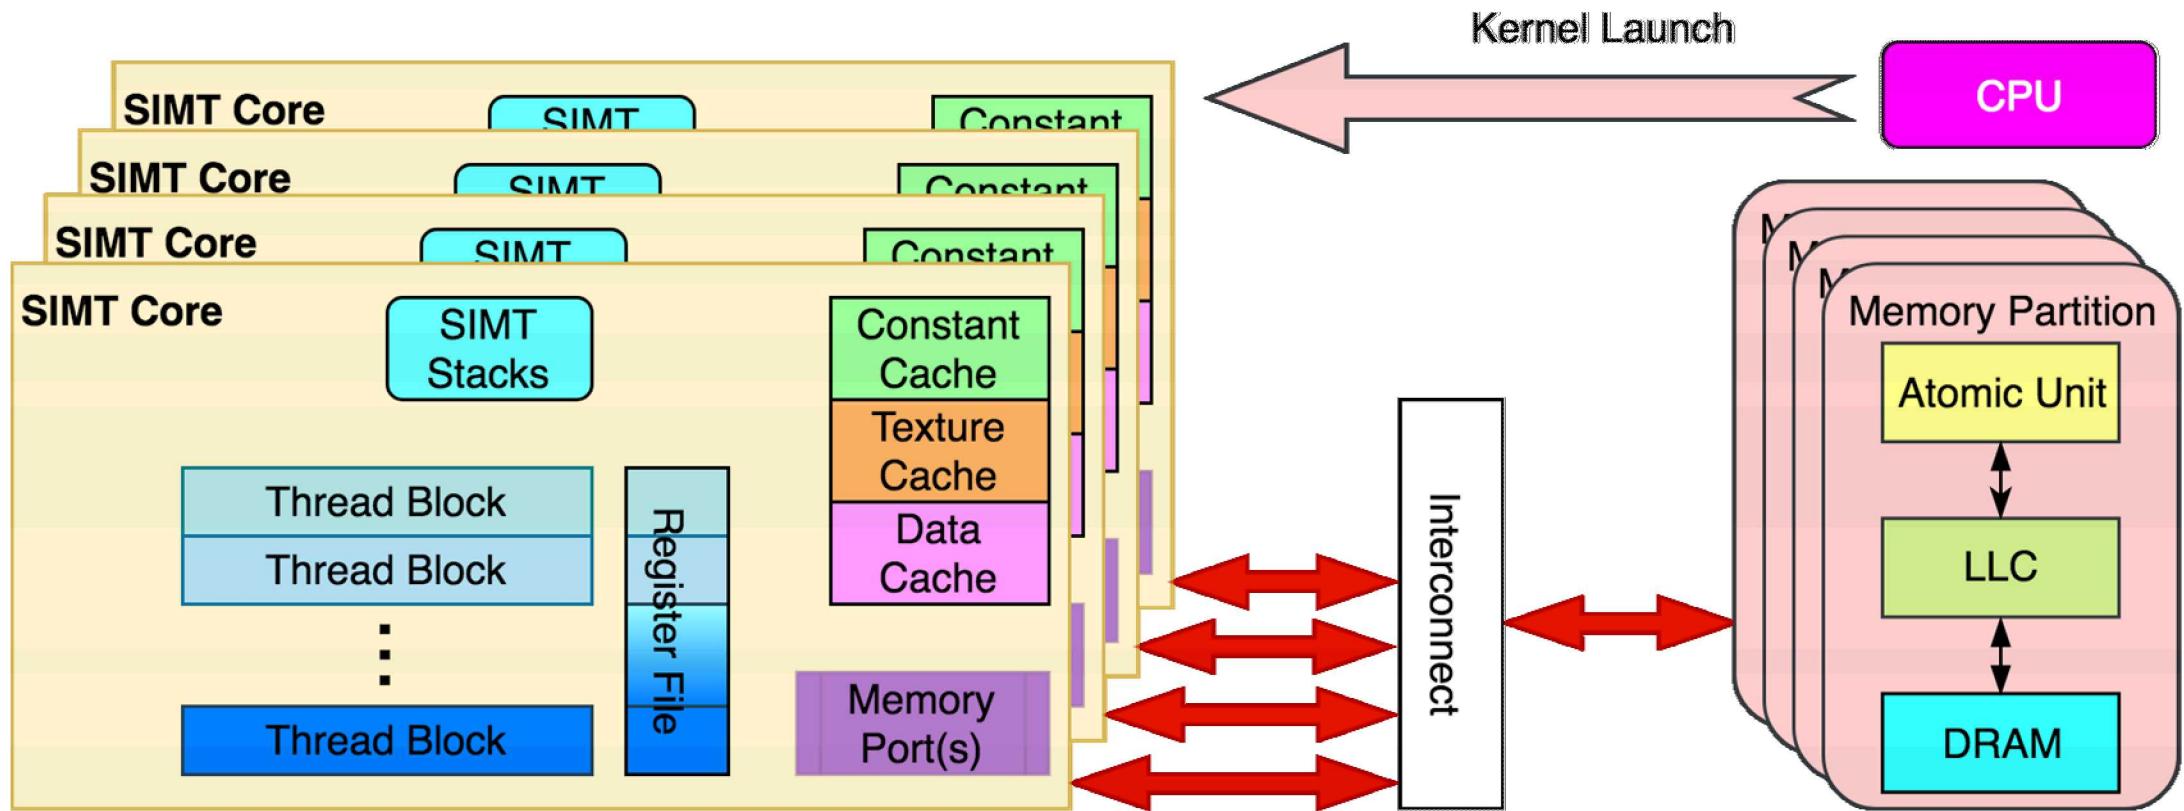
## Functional

- CTA
  - Pool of threads with access to shared memory
- Thread
  - PC + registers + allocated memory
- Instruction
  - 1:1 mapping with opcode and implementation function; simulated at issue

## Timing

- SIMT Core
  - Roughly equivalent to Fermi's SM
- Cache
  - L1 → Per-core, write-through/no-allocate
  - L2 → Device, write-back/no-allocate
- Interconnect
  - Intersim (Booksim)
- Backing store
- GDDR3/GDDR5

# GPGPUSim Integration



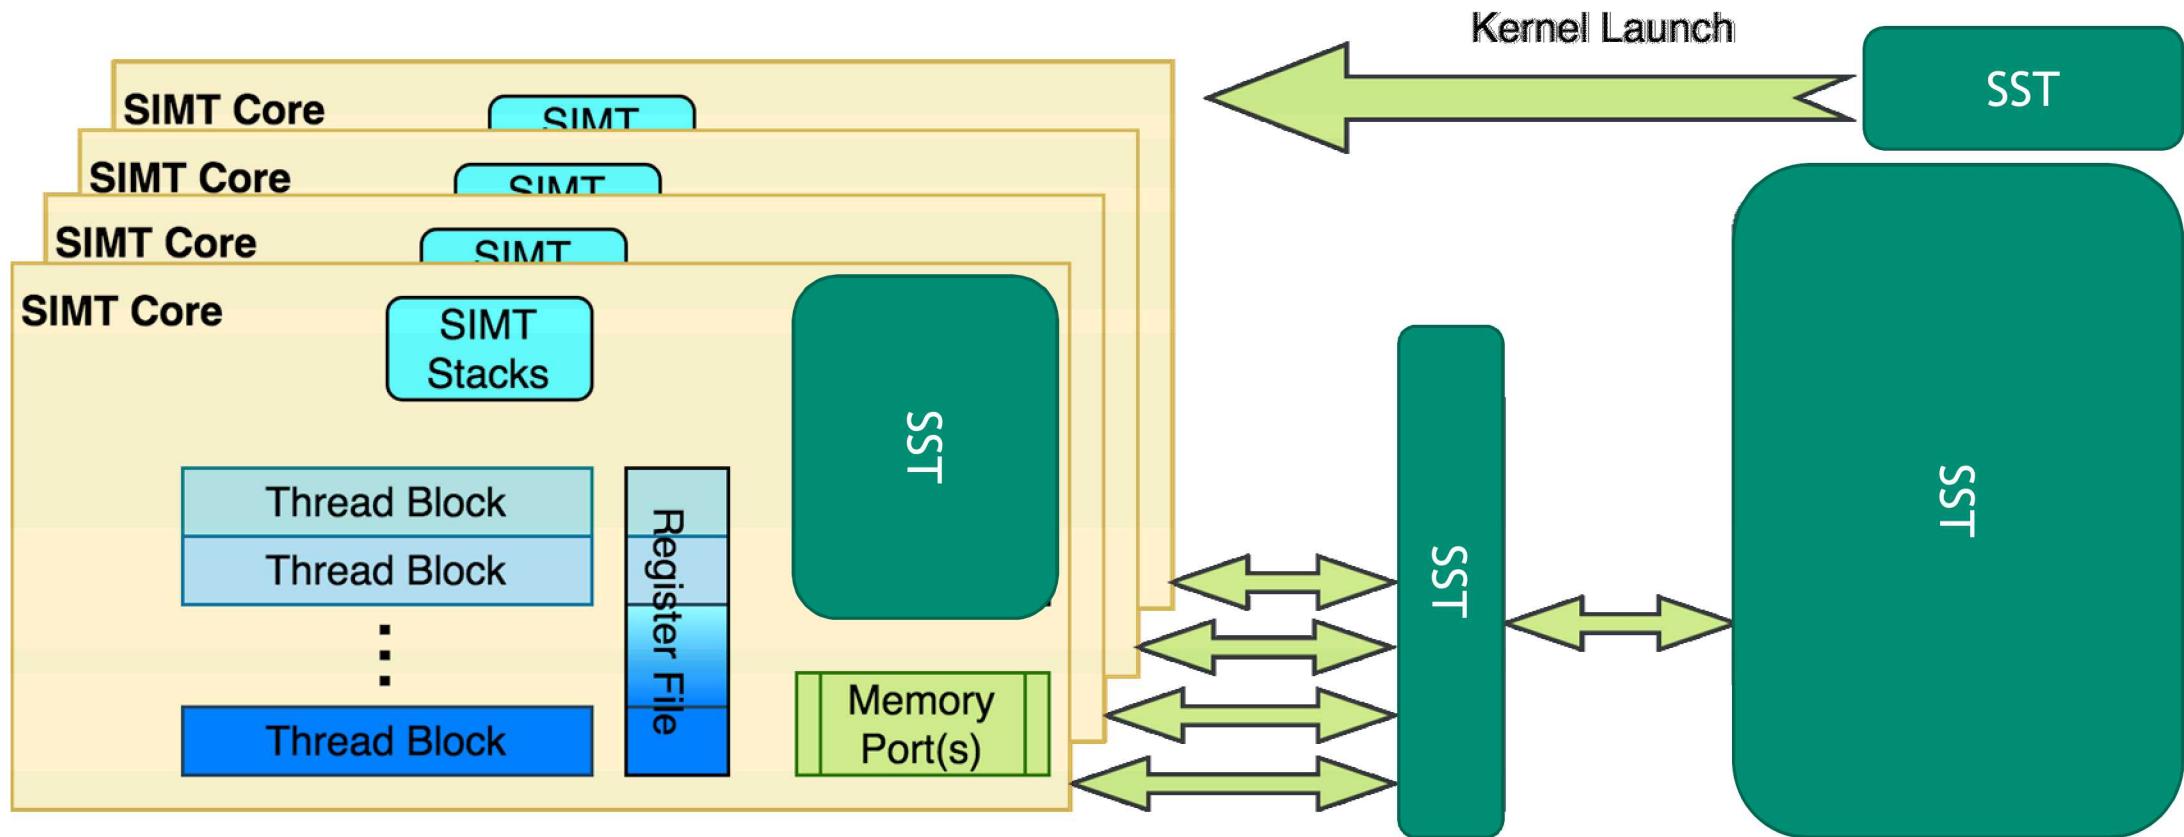
## Functional

- SIMT units track default functional model from GPGPUSim

## Timing

- SIMT units maintain execution timing; Booksim is used for the interconnect; and simple timing models are used for the memory partition

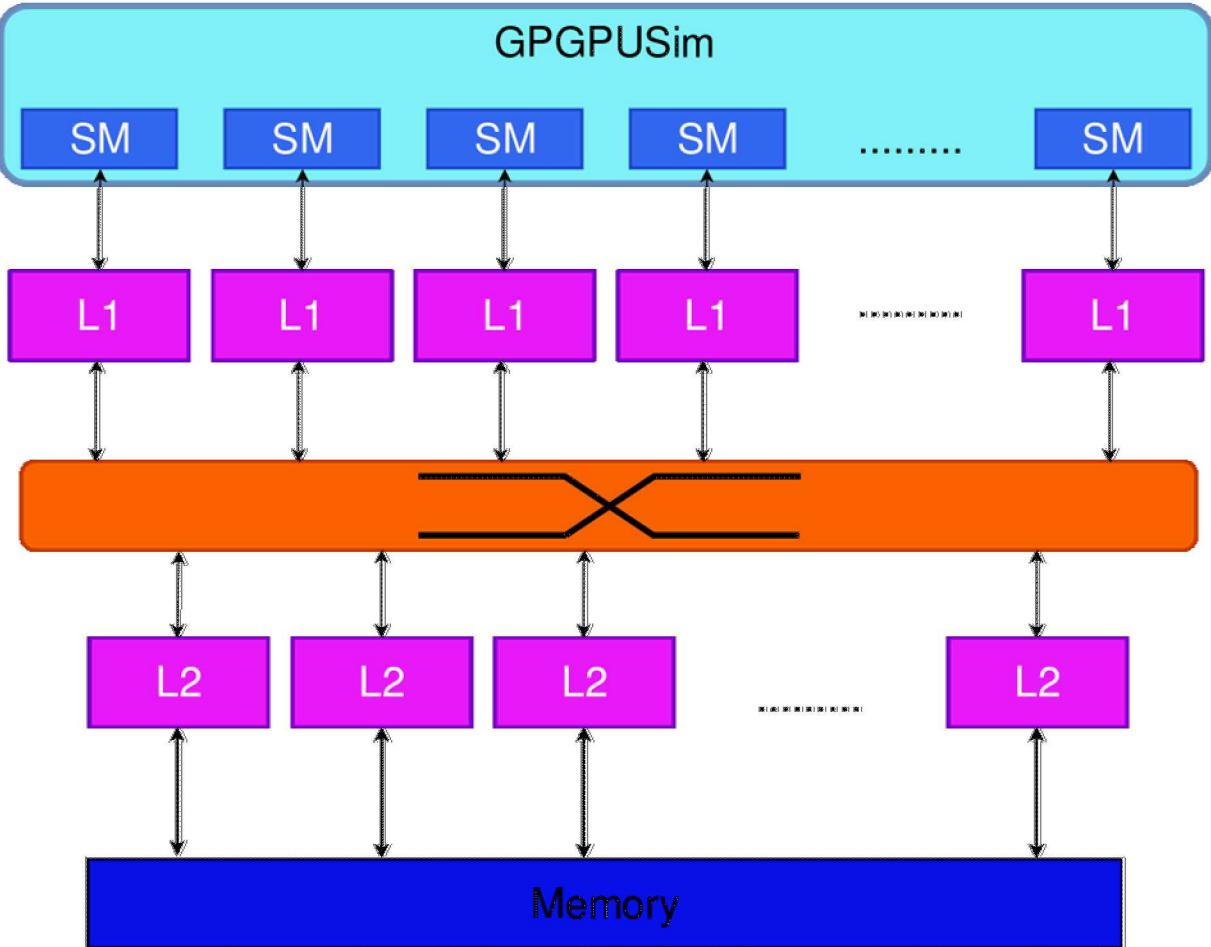
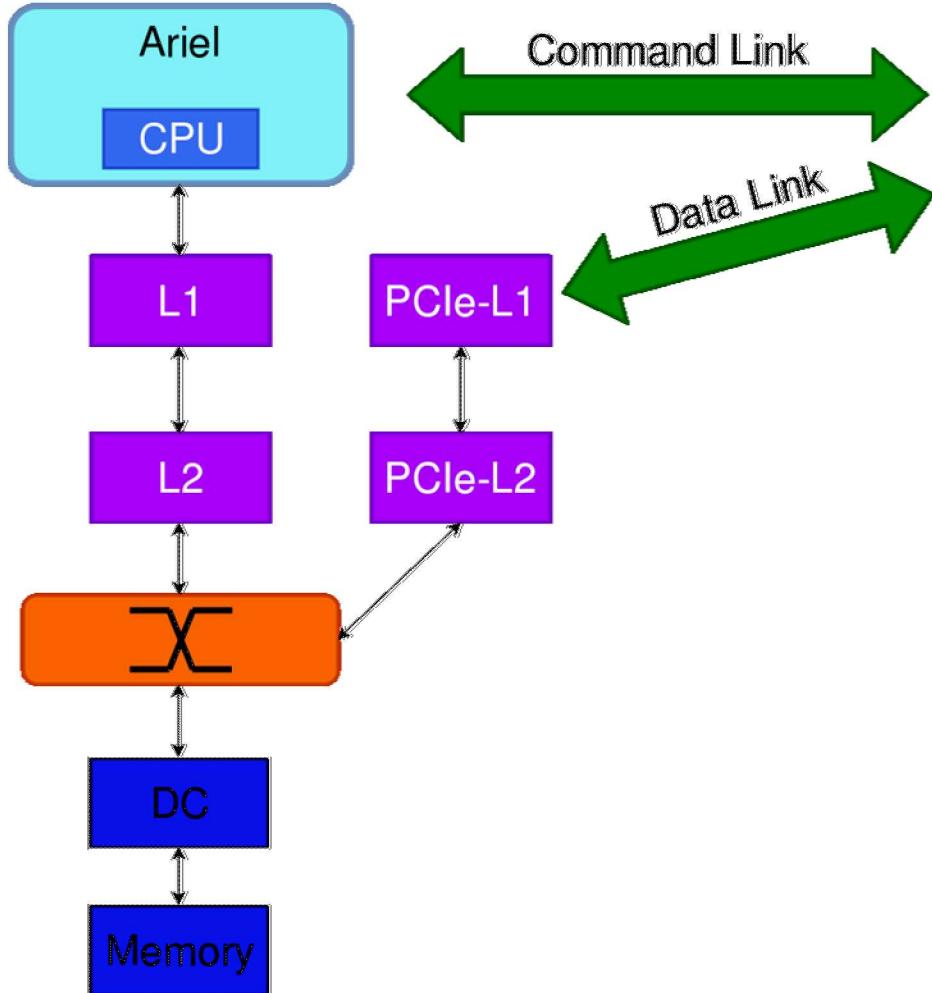
# GPGPUSim Integration



Plan to keep the SIMT units and replace everything else

- CPU → SST execution component (Ariel, Juno, etc.)
- Interconnect → SST networking component (Shogun, Merlin, Kingsley, etc.)
- Memory Partition & Caches → SST memHierarchy component for caches and various other backends for the backing store (DRAMSim, SimpleMem, Cramsim, etc.)

# GPGPUSim Integration



# Primary SST Components: Ariel

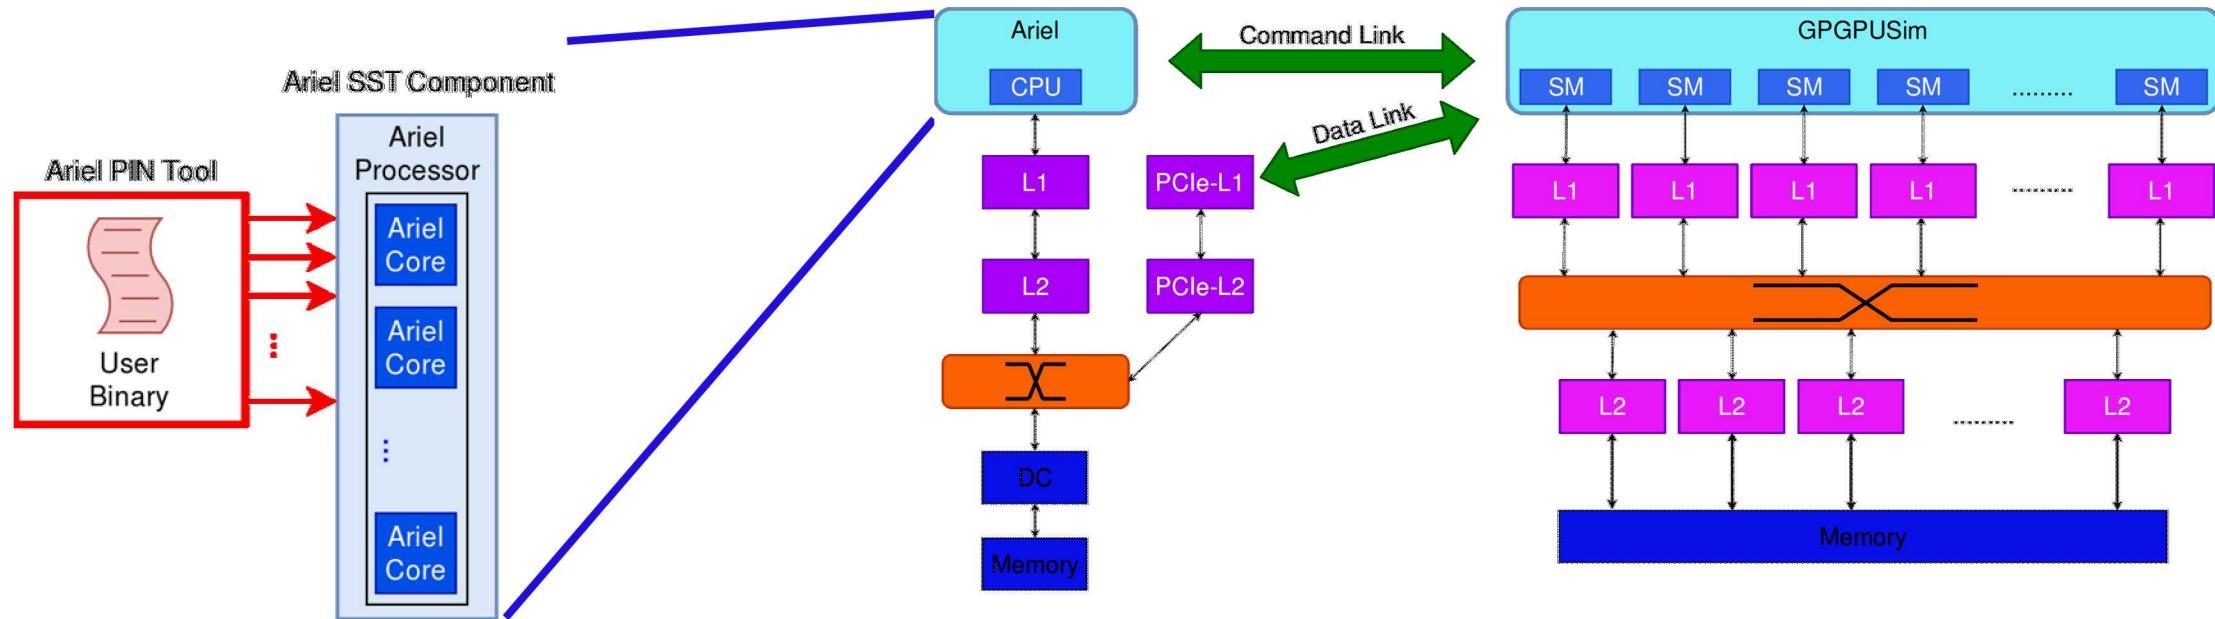
Frontend based on Intel PIN, which passes information to a SST component

- Supports OpenMP and MPI binaries

Faster than cycle-approximate models like Gem5 but slower than trace-based

Reasonable model of thread interactions

- Non-deterministic



# Primary SST Components: memHierarchy

Collection of interoperable memory system elements

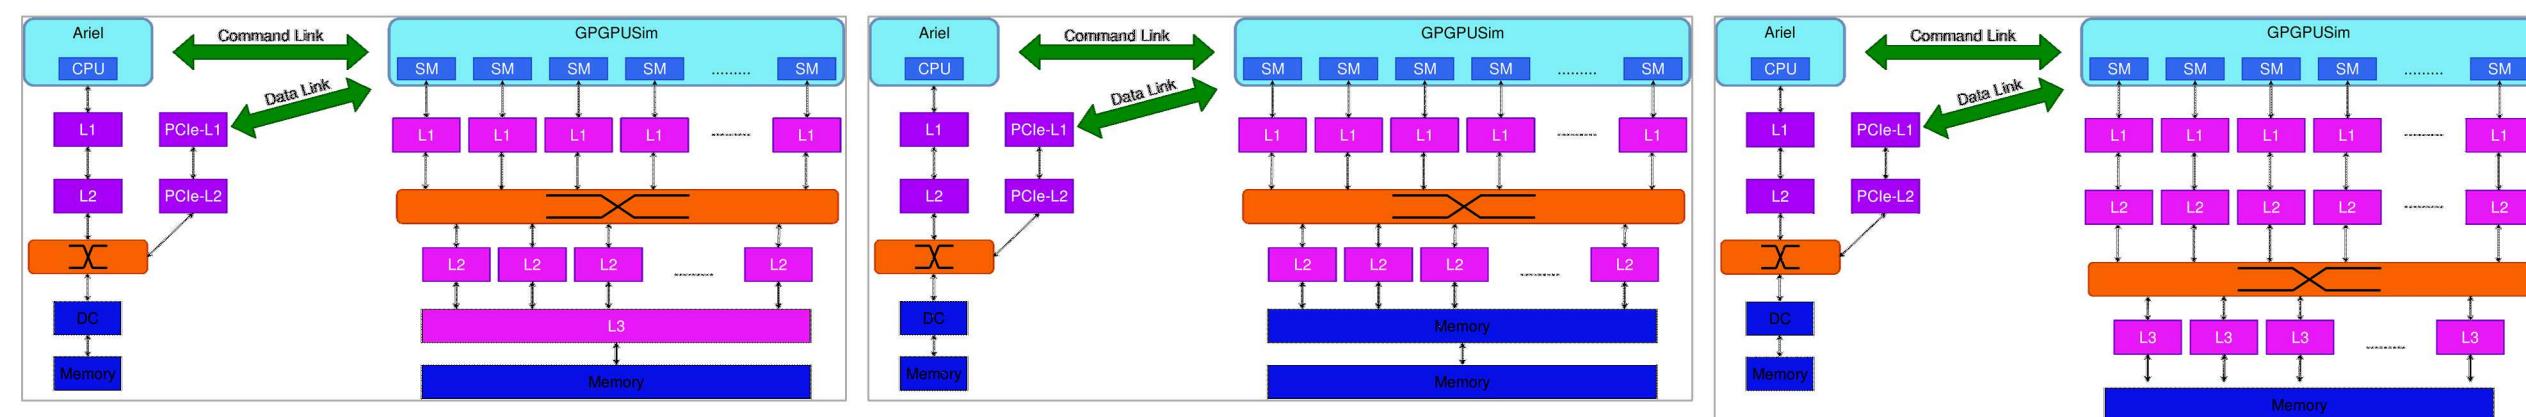
- Caches and directories
- Memory models (DDR, HBM, HMC, NVRAM, etc.)
- Memory controllers & network interfaces for memory (MemNICs)

Inter- and intra-socket coherence

Correlated with modern memory hierarchies

- HBM2/3 Evaluation on Many-core CPU for ECP
- Messier: A Detailed NVM-Based DIMM Model for the SST Simulation Framework

Will allow us to model any sort of memory arrangement that we can imagine!

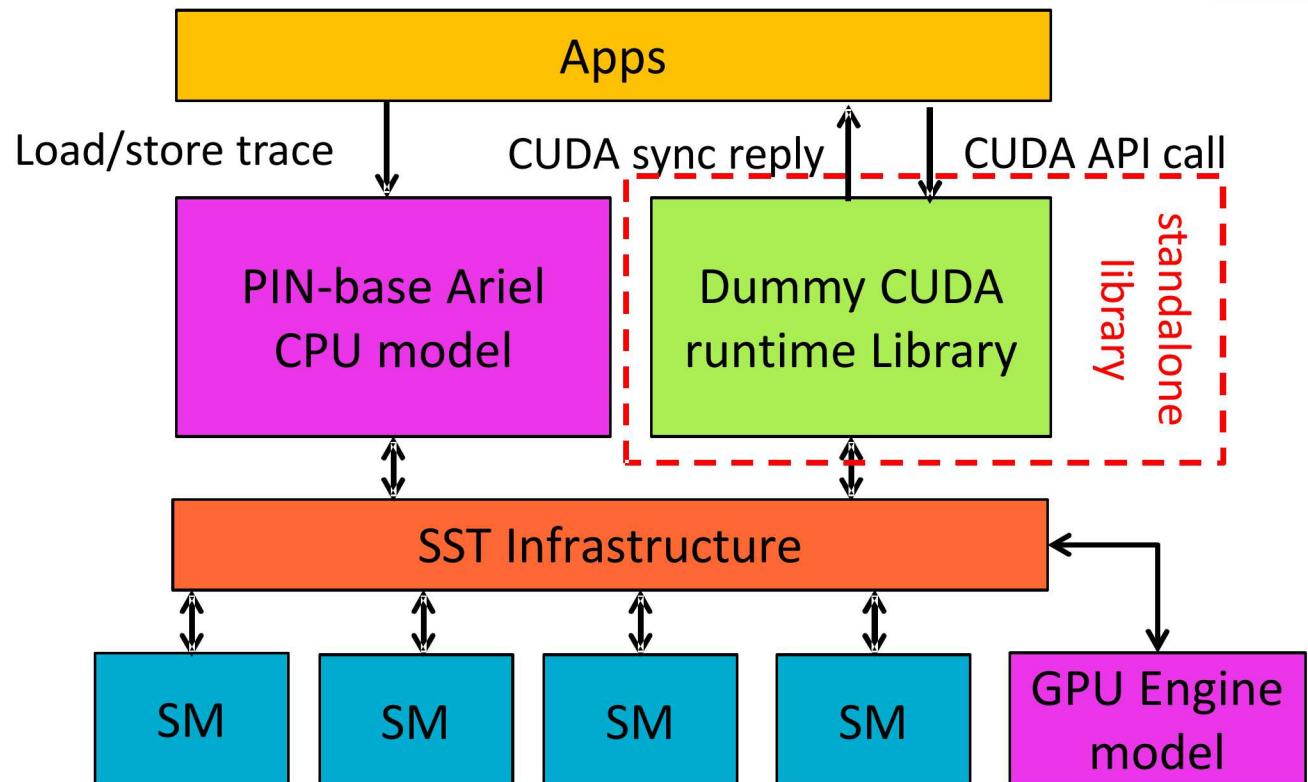


# CUDA API interception model

Standalone dummy CUDA lib which is independent of CPU model

Applications link dynamically with lib (shared object)

Dummy lib interacts with GPU engine for kernel launch and CUDA memory copies



API Calls Forwarded to GPU Model	
<code>__cudaRegisterFatBinary</code>	<code>cudaFree</code>
<code>__cudaRegisterFunction</code>	<code>cudaLaunch</code>
<code>cudaMalloc</code>	<code>cudaGetLastError</code>
<code>cudaMemcpy</code>	<code>cudaRegisterVar</code>
<code>cudaConfigureCall</code>	<code>cudaSetupArgument</code>
<code>cudaOccupancyMaxActiveBlocksPerMultiprocessorWithFlags</code>	



# Using The GPU Component

# Using the Component

Currently hosted in a separate repository within the SST project

To build SST elements with the GPU component (assumes core is built):

1. git clone https://github.com/sstsimulator/sst-elements.git
2. git clone https://github.com/sstsimulator/sst-gpgpusim.git \$SST\_ELEMENTS\_ROOT/src/sst/elements/Gpgpusim
3. . src/sst/elements/Gpgpusim/sst-gpgpusim/setup\_environment
4. cd src/sst/elements/Gpgpusim/sst-gpgpusim
5. make
6. cp --preserve=links \$GPGPUSIM\_ROOT/lib/\$GPGPUSIM\_CONFIG/libcudart\_mod.so  
\$SST\_ELEMENTS\_ROOT/src/sst/elements/Gpgpusim/
7. cd \$SST\_ELEMENTS\_HOME
8. ./autogen
9. ./configure --prefix=\$SST\_ELEMENTS\_HOME --with-sst-core=\$SST\_CORE\_HOME --with-pin=\$PIN\_HOME --with-cuda=\$CUDA\_ROOT --with-gpgpusim=\$GPGPUSIM\_ROOT
10. make

Applications must be built using the shared cuda runtime library: --cudart=shared

If an application requires cuDNN or cuBLAS, it must be linked against the static Cuda library: -cublas\_static, -lcudnn\_static

# Using the Component

Running a simulation requires the user to specify both a SST model and a GPGPU model

- SST model contains wiring instructions and component definitions
  - Processor model, memory hierarchy, interconnect, etc.
- GPGPU model contains SIMT core definition
  - PTX generation, number of registers, number of warps, number of functional units, etc.

## SST Component

```
comp_gpu0 = sst.Component("gpu0", "Gpgpusim.Gpgpusim")
comp_gpu0.addParams({
    "clock" : "1200MHz",
    "gpu_cores" : "84",
    "verbose" : "0"
})
```

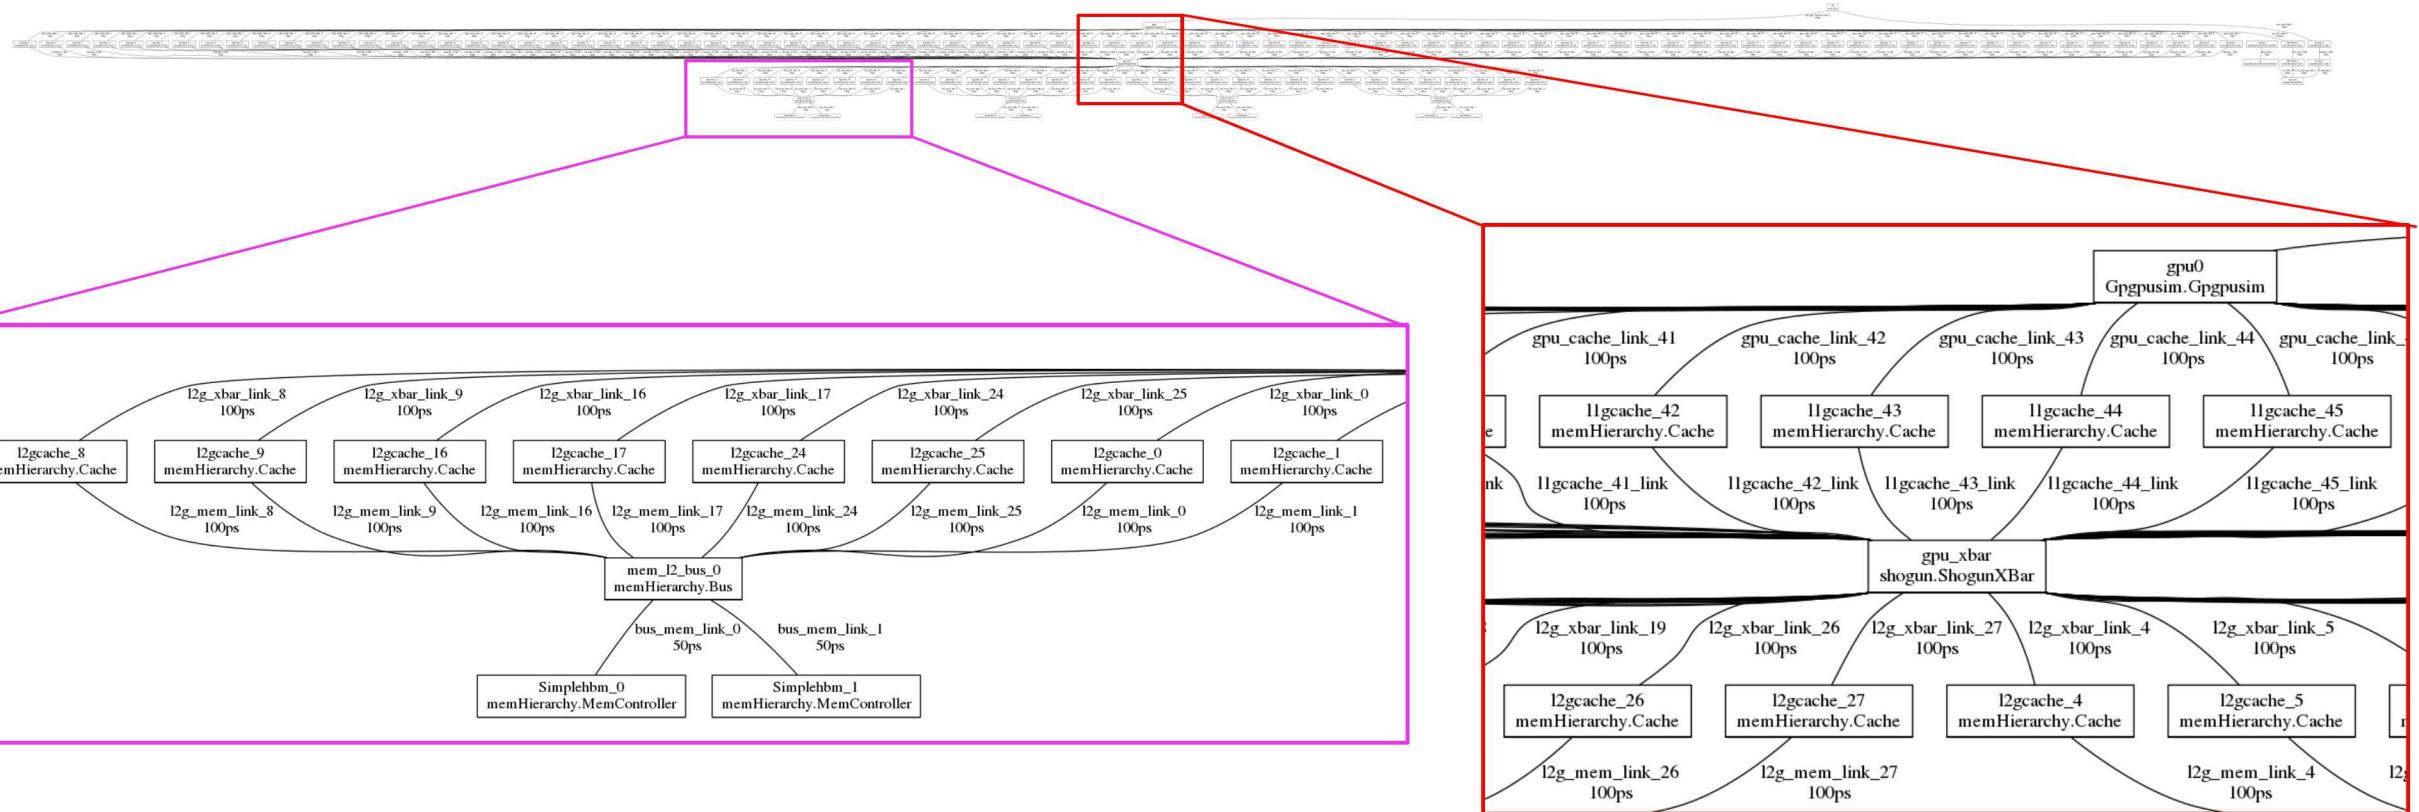
## GPGPU Model

```
-SST_mode 1
-gpgpu_ptx_force_max_capability 70
-gpgpu_clock_domains 1200.0:1200.0:1200.0:850.0
-gpgpu_shader_registers 65536
-gpgpu_occupancy_sm_number 70

-gpgpu_pipeline_widths 4,4,4,4,4,4,4,4,4,4,8,4,4
-gpgpu_num_sp_units 4
-gpgpu_num_sfu_units 4
-gpgpu_num_dp_units 4
-gpgpu_num_int_units 4
-gpgpu_tensor_core_avail 1
-gpgpu_num_tensor_core_units 4
```

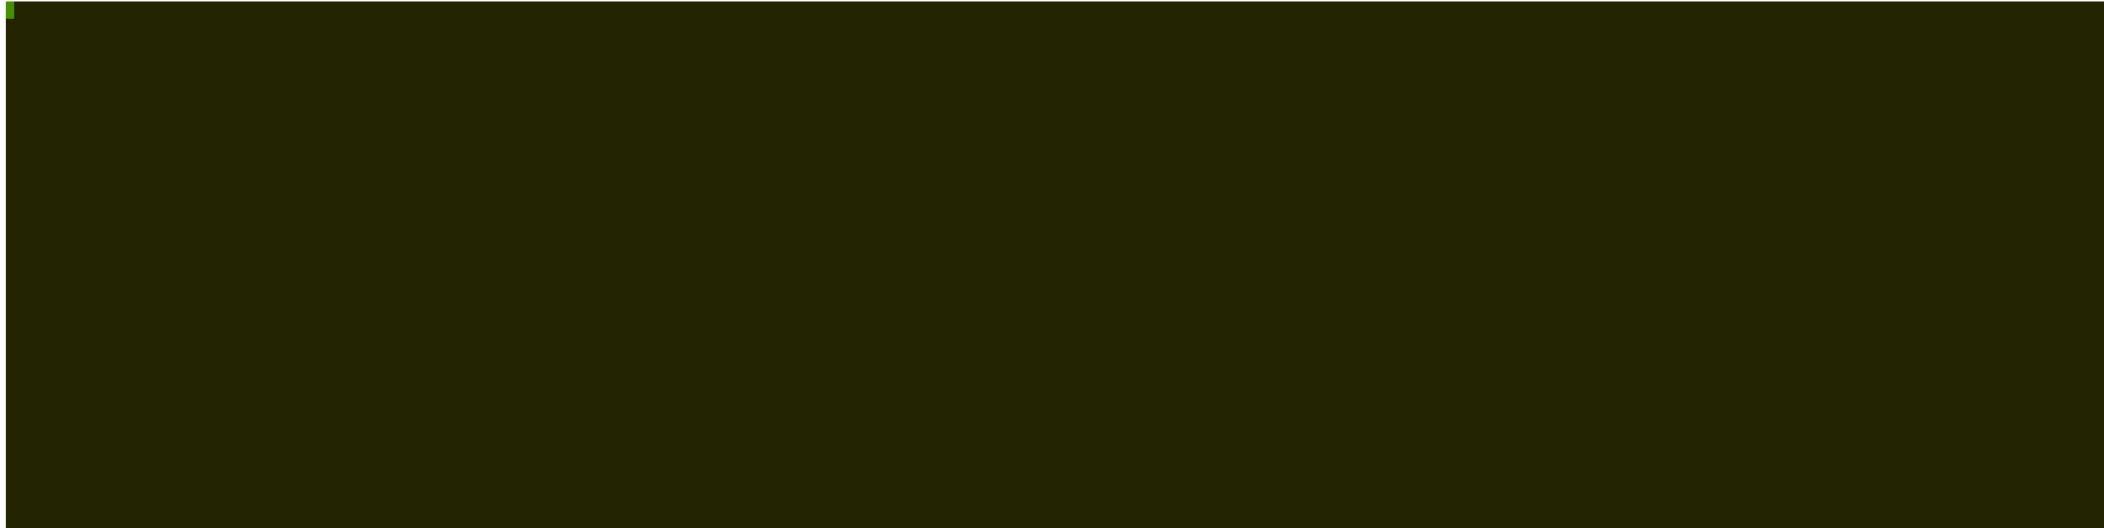
# Using the Component – vectorAdd

`sst --model-option="-c ariel-gpu-v100.cfg -s abs_kviTA.csv -x ./vectorAdd-163840" cuda-test.py`



# Using the Component – vectorAdd

```
sst --model-option="-c ariel-gpu-v100.cfg -s abs_kviTA.csv -x ./vectorAdd-163840" cuda-test.py
```



Statistics file for this model contains 19861 entries:

```
ComponentName, StatisticName, StatisticSubId, StatisticType, SimTime, Rank, Sum, SumSQ, Count, Min, Max
A0, read_requests, 0, Accumulator, 2423910320, 0, 10240, 10240, 10240, 1, 1
A0, write_requests, 0, Accumulator, 2423910320, 0, 20480, 20480, 20480, 1, 1
A0, read_request_sizes, 0, Accumulator, 2423910320, 0, 655360, 41943040, 10240, 64, 64
A0, write_request_sizes, 0, Accumulator, 2423910320, 0, 1310720, 83886080, 20480, 64, 64
A0, instruction_count, 0, Accumulator, 2423910320, 0, 0, 0, 0, 0, 0
A0, cycles, 0, Accumulator, 2423910320, 0, 6446570, 6446570, 6446570, 1, 1
A0, active_cycles, 0, Accumulator, 2423910320, 0, 226, 226, 226, 1, 1
...
...
```

# Performance Analysis

# Mean Error

	P9/V100	SST-GPGPU	Error
<b>vectorAdd</b>	5271	5751	9.093682
<b>lud</b>	494519	605685	22.47961
<b>lulesh</b>	12454750	11896477	4.48241

Complex model interactions make it difficult to pinpoint where models are lacking in detail or are incorrect

Scale of simulation – number of components and links present a challenge

- Full debug produces multi-TiB trace
- No errors, so have to isolate anomalies in tens of thousands of statistics

# Functional Testing – Kokkos Kernels

Kokkos-Kernels built with:

```
KOKKOSKERNELS_SCALARS=double  
KOKKOSKERNELS_LAYOUTS=LayoutLeft  
KOKKOSKERNELS_ORDINALS=int  
KOKKOSKERNELS_OFFSETS=int  
KOKKOS_DEVICES=Cuda,Serial  
KOKKOS_ARCHS=BDW,Pascal60
```

Tool	Version
gcc	4.8.5
CUDA SDK	8.0.61
Kokkos	2.8.00
Kokkos Kernels	2.8.00
sst-core	devel-d8042ec
sst-elements	devel-b3ba937
Pin	2.14.71313

Compiling for Broadwell and sm\_60 for validation testing on Doom

# Functional Testing – Kokkos Kernels Unit Test Results

1	abs_double		31	sparse_spadd_double_int_int_TestExecSpace		61	batched_scalar_serial_gemv_t_double_double
2	abs_mv_double		32	sparse_gauss_seidel_double_int_int_TestExecSpace		62	batched_scalar_serial_trsv_l_nt_u_double_double
3	asum_double		33	sparse_block_gauss_seidel_double_int_int_TestExecSpace		63	batched_scalar_serial_trsv_l_nt_n_double_double
4	axpby_double		34	sparse_crsmatrix_double_int_int_TestExecSpace		64	batched_scalar_serial_trsv_u_nt_u_double_double
5	axpby_mv_double		35	sparse_blkcrsmatrix_double_int_int_TestExecSpace		65	batched_scalar_serial_trsv_u_nt_n_double_double
6	axpy_double		36	sparse_replaceSumIntoLonger_double_int_int_TestExecSpace		66	batched_scalar_team_set_double_double
7	axpy_mv_double		37	sparse_replaceSumInto_double_int_int_TestExecSpace		67	batched_scalar_team_scale_double_double
8	dot_double		38	graph_graph_color_double_int_int_TestExecSpace		68	batched_scalar_team_gemm_nt_nt_double_double
9	dot_mv_double		39	graph_graph_color_deterministic_double_int_int_TestExecSpace		69	batched_scalar_team_gemm_t_nt_double_double
10	mult_double		40	graph_graph_color_d2_double_int_int_TestExecSpace		70	batched_scalar_team_gemm_nt_t_double_double
11	mult_mv_double		41	common_ArithTraits		71	batched_scalar_team_gemm_t_t_double_double
12	nrm1_double		42	common_set_bit_count		72	batched_scalar_team_trsm_ll_nt_u_double_double
13	nrm1_mv_double		43	common_ffs		73	batched_scalar_team_trsm_ll_nt_n_double_double
14	nrm2_double		44	batched_scalar_serial_set_double_double		74	batched_scalar_team_trsm_ll_u_nt_u_double_double
15	nrm2_mv_double		45	batched_scalar_serial_scale_double_double		75	batched_scalar_team_trsm_ll_u_nt_n_double_double
16	nrm2_squared_double		46	batched_scalar_serial_gemm_nt_nt_double_double		76	batched_scalar_team_trsm_r_u_nt_u_double_double
17	nrm2_squared_mv_double		47	batched_scalar_serial_gemm_t_nt_double_double		77	batched_scalar_team_trsm_r_u_nt_n_double_double
18	nrminf_double		48	batched_scalar_serial_gemm_nt_t_double_double		78	batched_scalar_team_trsm_ll_t_u_double_double
19	nrminf_mv_double		49	batched_scalar_serial_gemm_t_t_double_double		79	batched_scalar_team_trsm_ll_t_n_double_double
20	reciprocal_double		50	batched_scalar_serial_trsm_ll_nt_u_double_double		80	batched_scalar_team_trsm_ll_u_t_u_double_double
21	reciprocal_mv_double		51	batched_scalar_serial_trsm_ll_nt_n_double_double		81	batched_scalar_team_trsm_ll_u_t_n_double_double
22	scal_double		52	batched_scalar_serial_trsm_ll_u_nt_u_double_double		82	batched_scalar_team_gemv_nt_double_double
23	scal_mv_double		53	batched_scalar_serial_trsm_ll_u_nt_n_double_double		83	batched_scalar_team_gemv_t_double_double
24	sum_double		54	batched_scalar_serial_trsm_r_u_nt_u_double_double		84	batched_scalar_serial_lu_double
25	sum_mv_double		55	batched_scalar_serial_trsm_r_u_nt_n_double_double		85	batched_scalar_serial_inverselu_double
26	update_double		56	batched_scalar_serial_trsm_ll_t_u_double_double		86	batched_scalar_serial_solvelu_double
27	update_mv_double		57	batched_scalar_serial_trsm_ll_t_n_double_double		87	batched_scalar_team_lu_double
28	gemv_double		58	batched_scalar_serial_trsm_ll_u_t_u_double_double		88	batched_scalar_team_inverselu_double
29	gemm_double		59	batched_scalar_serial_trsm_ll_u_t_n_double_double		89	batched_scalar_team_solvelu_double
30	sparse_spgemm_double_int_int_TestExecSpace		60	batched_scalar_serial_gemv_nt_double_double			

Currently passing 54/89 tests (60.67%)

# Functional Testing – Kokkos Kernels Unit Test Results

1	abs_double		31	sparse_spadd_double_int_int_TestExecSpace		61	batched_scalar_serial_gemv_t_double_double
2	abs_mv_double		32	sparse_gauss_seidel_double_int_int_TestExecSpace		62	batched_scalar_serial_trsv_l_nt_u_double_double
3	asum_double		33	sparse_block_gauss_seidel_double_int_int_TestExecSpace		63	batched_scalar_serial_trsv_l_nt_n_double_double
4	axpby_double		34	sparse_crsmatrix_double_int_int_TestExecSpace		64	batched_scalar_serial_trsv_u_nt_u_double_double
5	axpby_mv_double		35	sparse_blkcrsmatrix_double_int_int_TestExecSpace		65	batched_scalar_serial_trsv_u_nt_n_double_double
6	axpy_double		36	sparse_replaceSumIntoLonger_double_int_int_TestExecSpace		66	batched_scalar_team_set_double_double
7	axpy_mv_double		37	sparse_replaceSumInto_double_int_int_TestExecSpace		67	batched_scalar_team_scale_double_double
8	dot_double		38	graph_graph_color_double_int_int_TestExecSpace		68	batched_scalar_team_gemm_nt_nt_double_double
9	dot_mv_double		39	graph_graph_color_deterministic_double_int_int_TestExecSpace		69	batched_scalar_team_gemm_t_nt_double_double
10	mult_double		40	graph_graph_color_d2_double_int_int_TestExecSpace		70	batched_scalar_team_gemm_nt_t_double_double
11	mult_mv_double		41	common_ArithTraits		71	batched_scalar_team_gemm_t_t_double_double
12	nrm1_double		42	common_set_bit_count		72	batched_scalar_team_trsm_ll_nt_u_double_double
13	nrm1_mv_double		43	common_ffs		73	batched_scalar_team_trsm_ll_nt_n_double_double
14	nrm2_double		44	batched_scalar_serial_set_double_double		74	batched_scalar_team_trsm_ll_u_nt_u_double_double
15	nrm2_mv_double		45	batched_scalar_serial_scale_double_double		75	batched_scalar_team_trsm_ll_u_nt_n_double_double
16	nrm2_squared_double		46	batched_scalar_serial_gemm_nt_nt_double_double		76	batched_scalar_team_trsm_r_u_nt_u_double_double
17	nrm2_squared_mv_double		47	batched_scalar_serial_gemm_t_nt_double_double		77	batched_scalar_team_trsm_r_u_nt_n_double_double
18	nrminf_double		48	batched_scalar_serial_gemm_nt_t_double_double		78	batched_scalar_team_trsm_ll_t_u_double_double
19	nrminf_mv_double		49	batched_scalar_serial_gemm_t_t_double_double		79	batched_scalar_team_trsm_ll_t_n_double_double
20	reciprocal_double		50	batched_scalar_serial_trsm_ll_nt_u_double_double		80	batched_scalar_team_trsm_ll_u_t_u_double_double
21	reciprocal_mv_double		51	batched_scalar_serial_trsm_ll_nt_n_double_double		81	batched_scalar_team_trsm_ll_u_t_n_double_double
22	scal_double		52	batched_scalar_serial_trsm_ll_u_nt_u_double_double		82	batched_scalar_team_gemv_nt_double_double
23	scal_mv_double		53	batched_scalar_serial_trsm_ll_u_nt_n_double_double		83	batched_scalar_team_gemv_t_double_double
24	sum_double		54	batched_scalar_serial_trsm_r_u_nt_u_double_double		84	batched_scalar_serial_lu_double
25	sum_mv_double		55	batched_scalar_serial_trsm_r_u_nt_n_double_double		85	batched_scalar_serial_inverselu_double
26	update_double		56	batched_scalar_serial_trsm_ll_t_u_double_double		86	batched_scalar_serial_solvelu_double
27	update_mv_double		57	batched_scalar_serial_trsm_ll_t_n_double_double		87	batched_scalar_team_lu_double
28	gemv_double		58	batched_scalar_serial_trsm_ll_u_t_u_double_double		88	batched_scalar_team_inverselu_double
29	gemm_double		59	batched_scalar_serial_trsm_ll_u_t_n_double_double		89	batched_scalar_team_solvelu_double
30	sparse_spgemm_double_int_int_TestExecSpace		60	batched_scalar_serial_gemv_nt_double_double			

Tests in pink fail because the parser cannot locate a post-dominator in some kernels; this bug is also present in the public branch of GPGPUSSim

# Functional Testing – Kokkos Kernels Unit Test Results

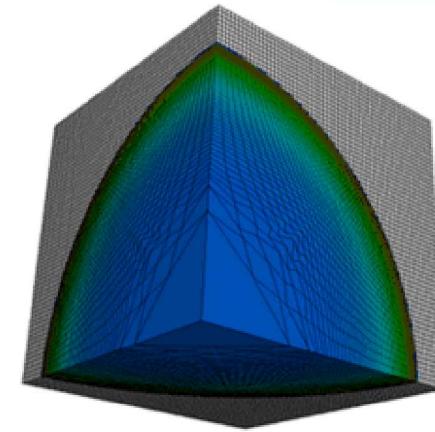
1	abs_double		31	sparse_spadd_double_int_int_TestExecSpace		61	batched_scalar_serial_gemv_t_double_double
2	abs_mv_double		32	sparse_gauss_seidel_double_int_int_TestExecSpace		62	batched_scalar_serial_trsv_l_nt_u_double_double
3	asum_double		33	sparse_block_gauss_seidel_double_int_int_TestExecSpace		63	batched_scalar_serial_trsv_l_nt_n_double_double
4	axpby_double		34	sparse_crsmatrix_double_int_int_TestExecSpace		64	batched_scalar_serial_trsv_u_nt_u_double_double
5	axpby_mv_double		35	sparse_blkcrsmatrix_double_int_int_TestExecSpace		65	batched_scalar_serial_trsv_u_nt_n_double_double
6	axpy_double		36	sparse_replaceSumIntoLonger_double_int_int_TestExecSpace		66	batched_scalar_team_set_double_double
7	axpy_mv_double		37	sparse_replaceSumInto_double_int_int_TestExecSpace		67	batched_scalar_team_scale_double_double
8	dot_double		38	graph_graph_color_double_int_int_TestExecSpace		68	batched_scalar_team_gemm_nt_nt_double_double
9	dot_mv_double		39	graph_graph_color_deterministic_double_int_int_TestExecSpace		69	batched_scalar_team_gemm_t_nt_double_double
10	mult_double		40	graph_graph_color_d2_double_int_int_TestExecSpace		70	batched_scalar_team_gemm_nt_t_double_double
11	mult_mv_double		41	common_ArithTraits		71	batched_scalar_team_gemm_t_t_double_double
12	nrm1_double		42	common_set_bit_count		72	batched_scalar_team_trsm_l_l_nt_u_double_double
13	nrm1_mv_double		43	common_ffs		73	batched_scalar_team_trsm_l_l_nt_n_double_double
14	nrm2_double		44	batched_scalar_serial_set_double_double		74	batched_scalar_team_trsm_l_u_nt_u_double_double
15	nrm2_mv_double		45	batched_scalar_serial_scale_double_double		75	batched_scalar_team_trsm_l_u_nt_n_double_double
16	nrm2_squared_double		46	batched_scalar_serial_gemm_nt_nt_double_double		76	batched_scalar_team_trsm_r_u_nt_u_double_double
17	nrm2_squared_mv_double		47	batched_scalar_serial_gemm_t_nt_double_double		77	batched_scalar_team_trsm_r_u_nt_n_double_double
18	nrminf_double		48	batched_scalar_serial_gemm_nt_t_double_double		78	batched_scalar_team_trsm_l_l_t_u_double_double
19	nrminf_mv_double		49	batched_scalar_serial_gemm_t_t_double_double		79	batched_scalar_team_trsm_l_l_t_n_double_double
20	reciprocal_double		50	batched_scalar_serial_trsm_l_l_nt_u_double_double		80	batched_scalar_team_trsm_l_u_t_u_double_double
21	reciprocal_mv_double		51	batched_scalar_serial_trsm_l_l_nt_n_double_double		81	batched_scalar_team_trsm_l_u_t_n_double_double
22	scal_double		52	batched_scalar_serial_trsm_l_u_nt_u_double_double		82	batched_scalar_team_gemv_nt_double_double
23	scal_mv_double		53	batched_scalar_serial_trsm_l_u_nt_n_double_double		83	batched_scalar_team_gemv_t_double_double
24	sum_double		54	batched_scalar_serial_trsm_r_u_nt_u_double_double		84	batched_scalar_serial_lu_double
25	sum_mv_double		55	batched_scalar_serial_trsm_r_u_nt_n_double_double		85	batched_scalar_serial_inverselu_double
26	update_double		56	batched_scalar_serial_trsm_l_l_t_u_double_double		86	batched_scalar_serial_solvelu_double
27	update_mv_double		57	batched_scalar_serial_trsm_l_l_t_n_double_double		87	batched_scalar_team_lu_double
28	gemv_double		58	batched_scalar_serial_trsm_l_u_t_u_double_double		88	batched_scalar_team_inverselu_double
29	gemm_double		59	batched_scalar_serial_trsm_l_u_t_n_double_double		89	batched_scalar_team_solvelu_double
30	sparse_spgemm_double_int_int_TestExecSpace		60	batched_scalar_serial_gemv_nt_double_double			

Tests in purple fail because of a bug(s) in the functional model that neither the SST nor the GPGPUSim team has been able to isolate

# Performance Analysis – Lulesh

One of the most widely used DOE mini-applications

- Developed by Lawrence Livermore National Laboratory
- Represents challenging hydrodynamics algorithms over unstructured meshes
  - Common in many high-performance computing centers and are particularly prevalent within the NNSA laboratories
  - Routinely counted in the top ten application codes in terms of CPU hours utilized



Lulesh was compiled with:

Tool	Version
gcc	4.8.5
CUDA SDK	9.1.85
Lulesh	1.0
sst-core	devel-d8042ec
sst-elements	devel-b3ba937
Pin	2.14.71313

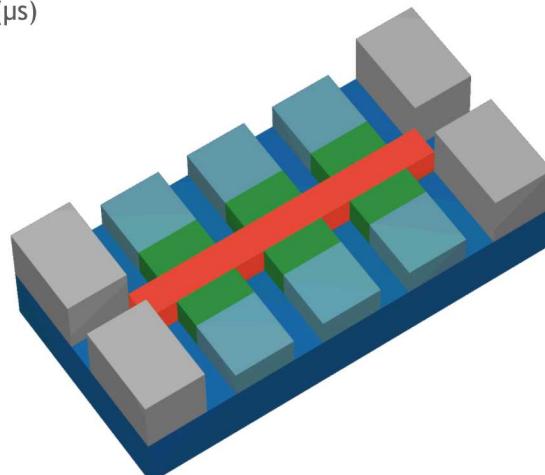
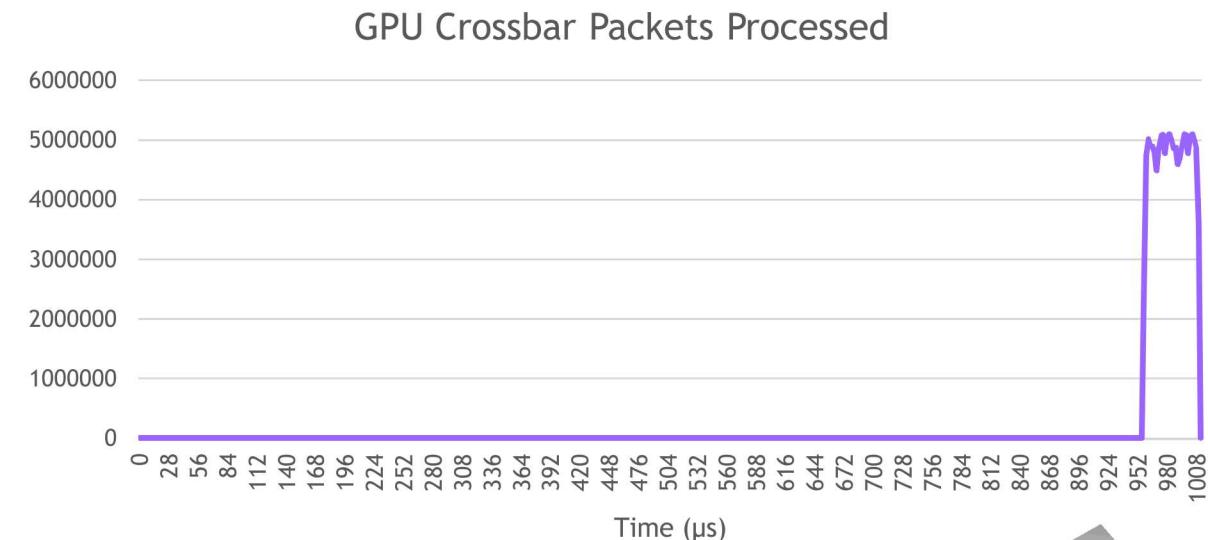
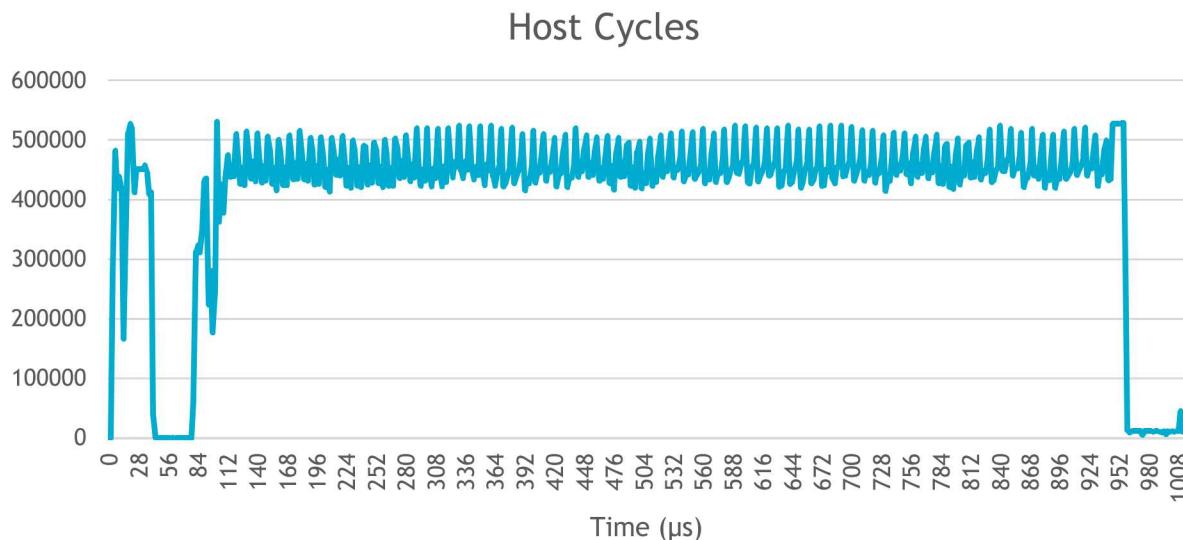
Lulesh was run with:

- $s=22$  ( $22 \times 22 \times 22$  elements)
- Iterations=50

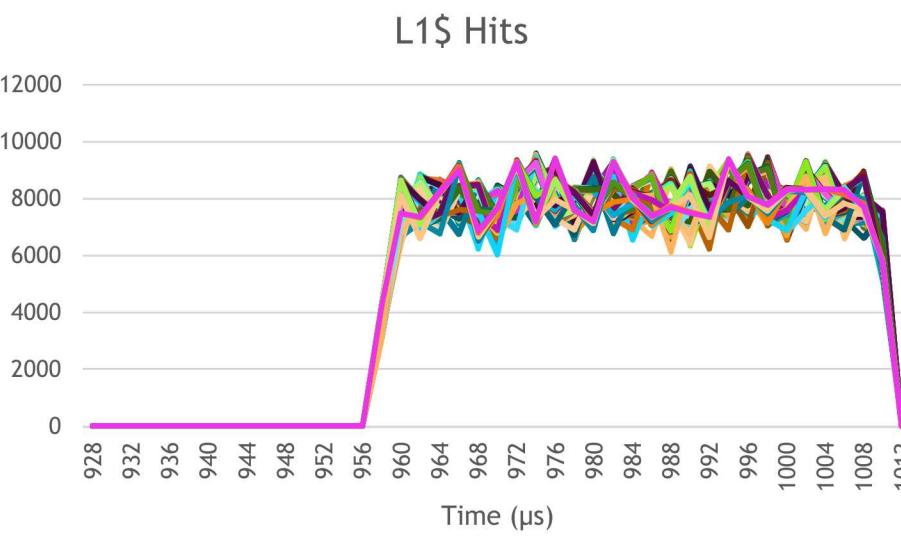
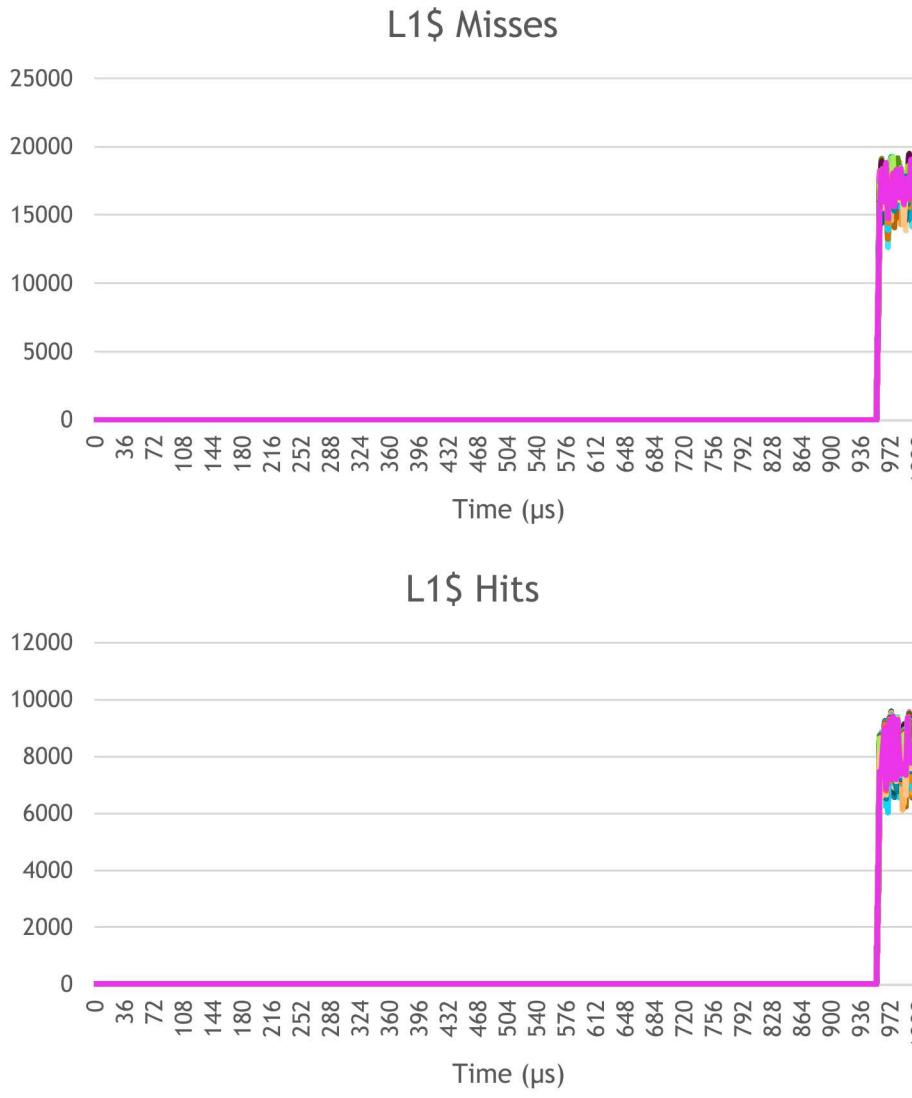
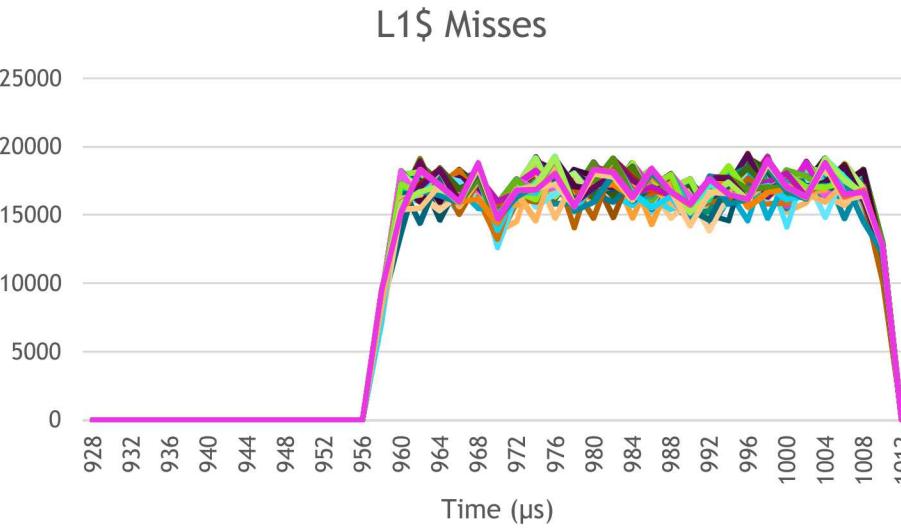
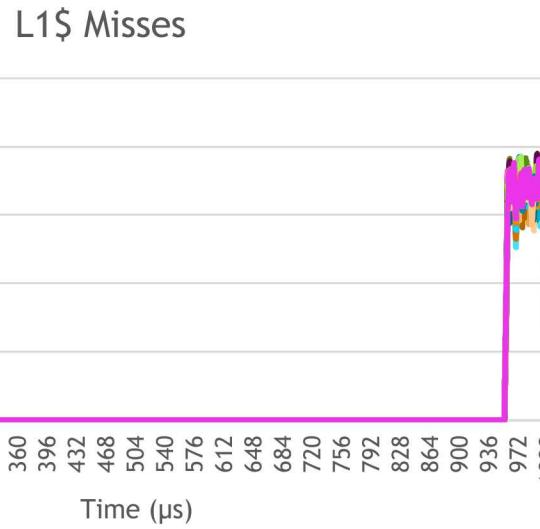
# Performance Analysis – Lulesh

SST is capable of providing periodic statistic dumps for all of the currently loaded components

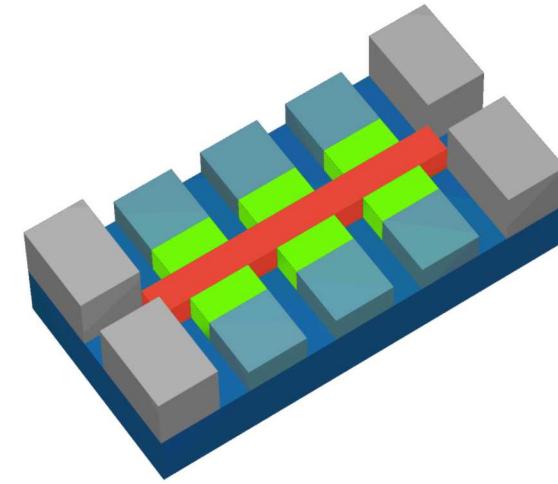
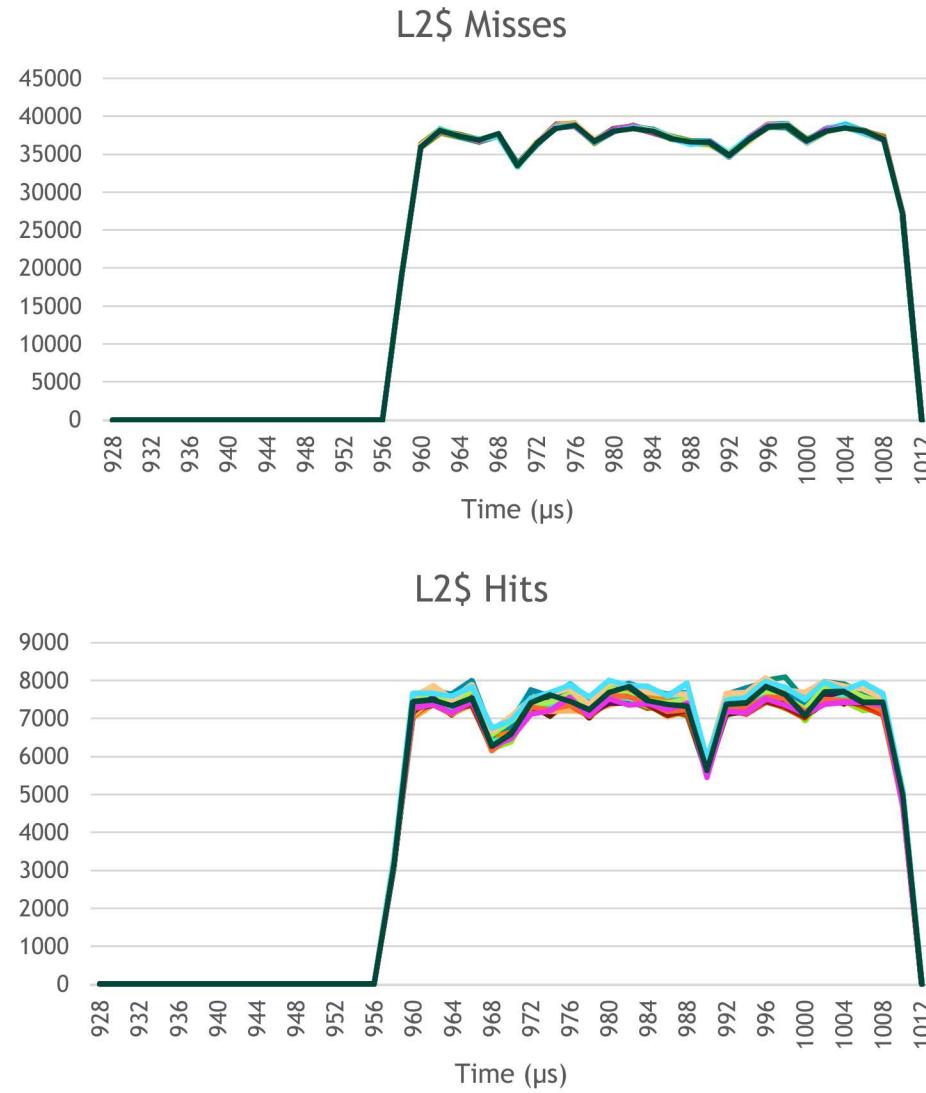
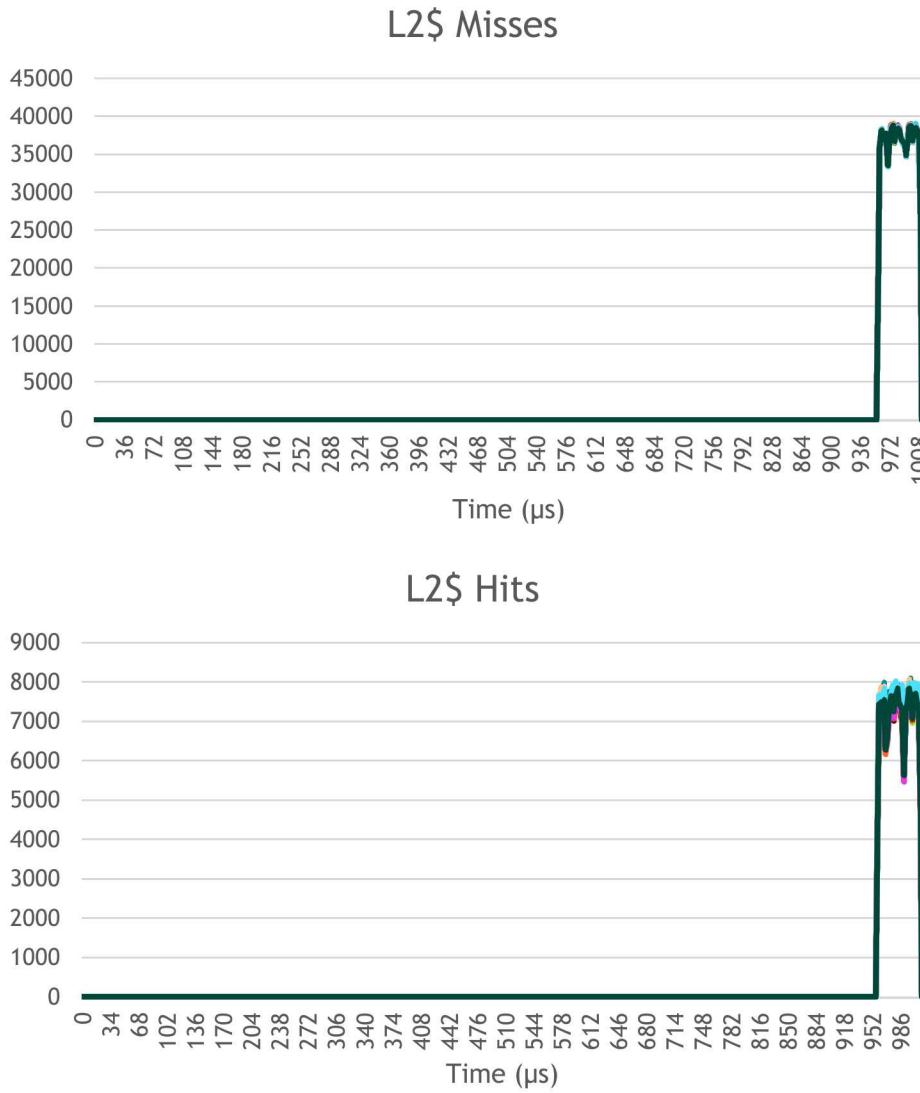
- Valuable for fine-grain performance analysis of applications



# Performance Analysis L1 Cache – Lulesh

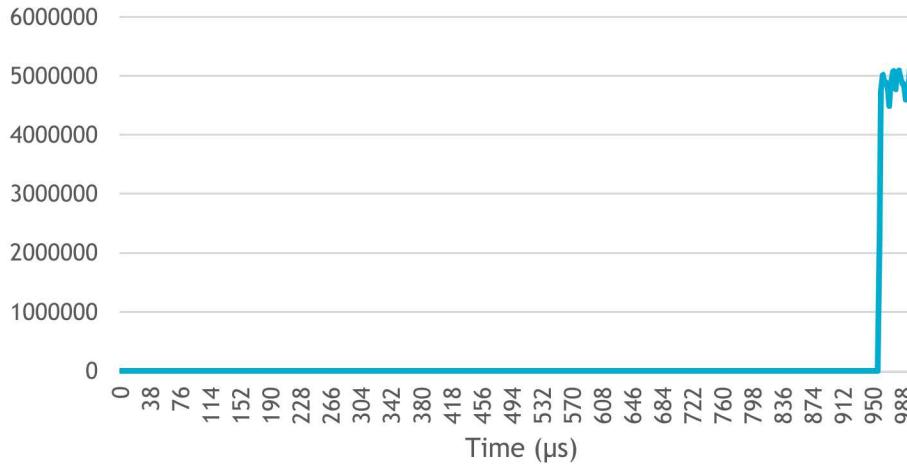


# Performance Analysis L2 Cache – Lulesh

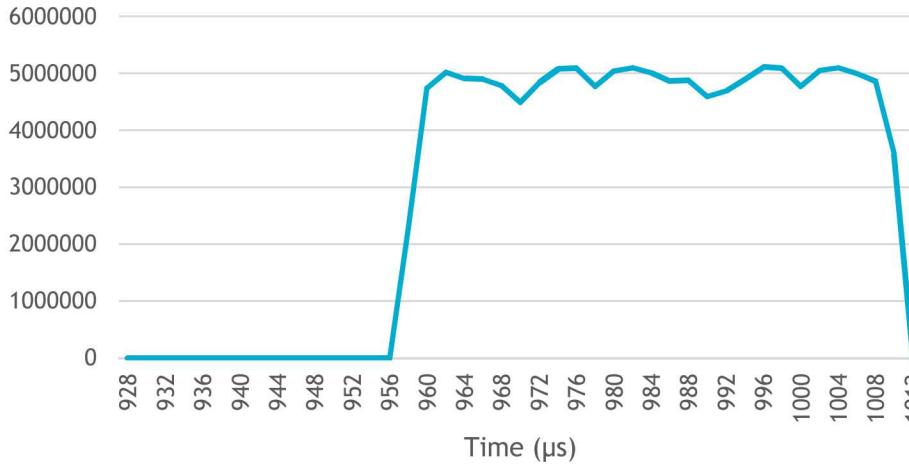


# Performance Analysis Crossbar – Lulesh

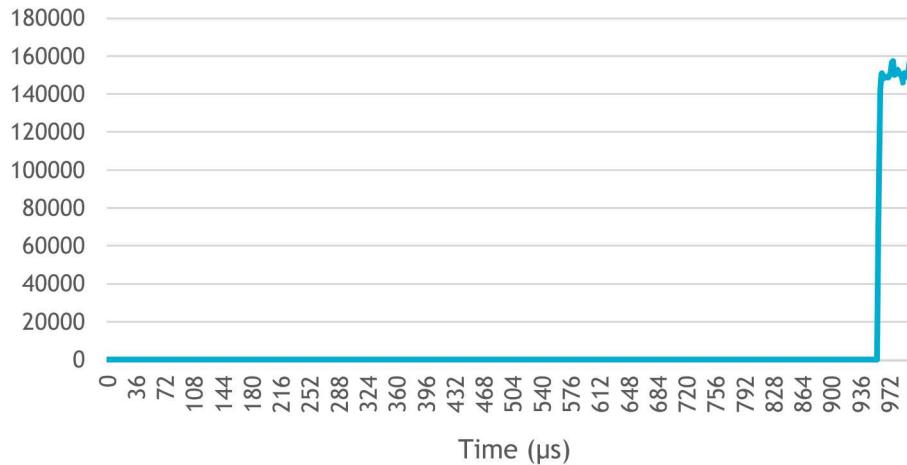
Crossbar Packets Processed



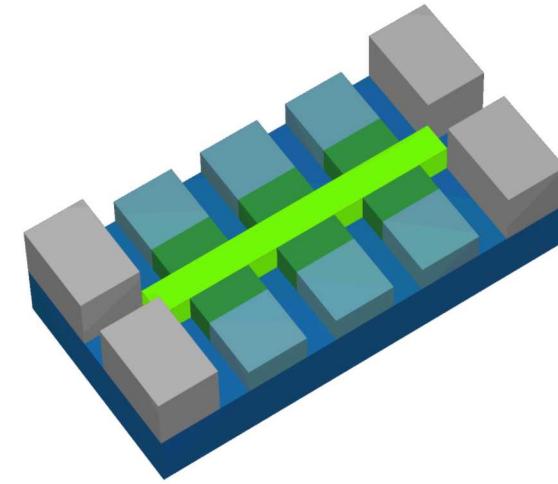
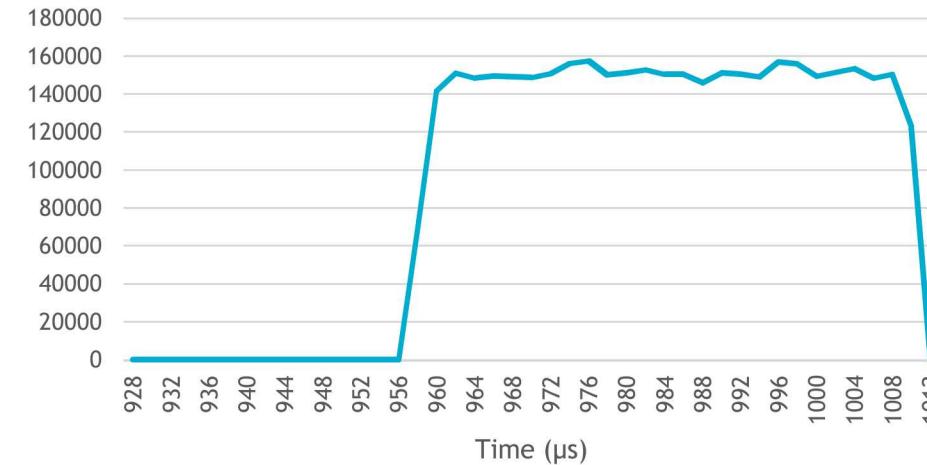
Crossbar Packets Processed



Crossbar Cycles Active



Crossbar Cycles Active



# Performance Analysis Memory Controller– Lulesh



27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

27

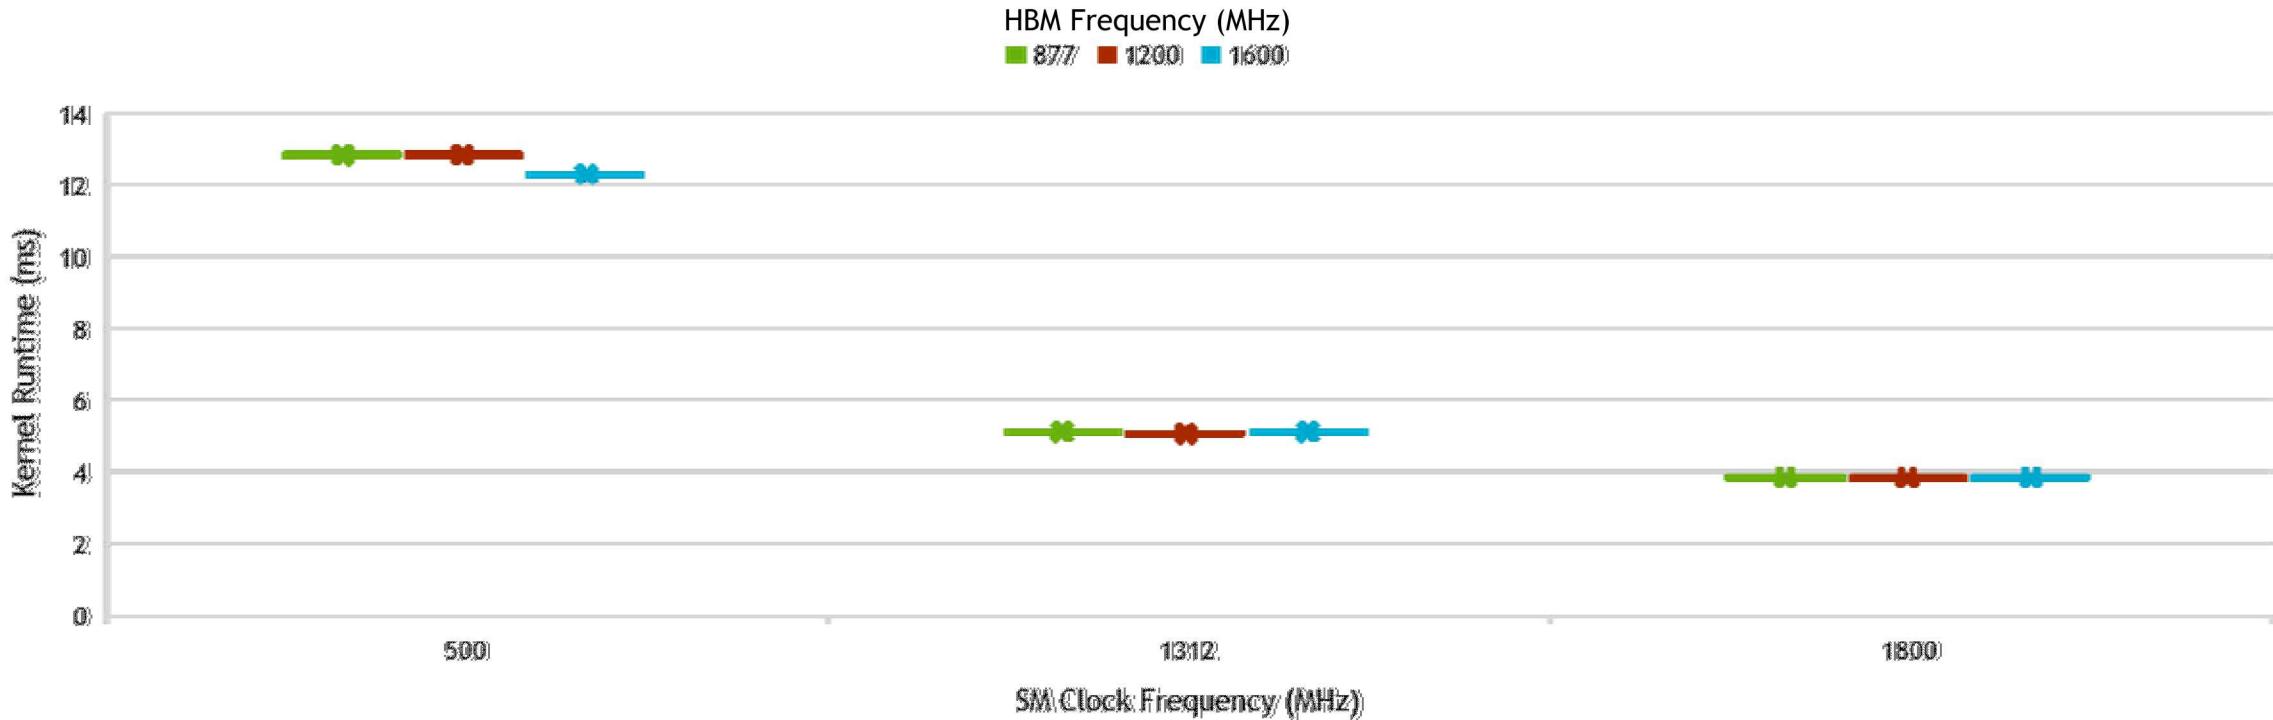
27

27

27

27&lt;/

# Design Space Exploration – Lulesh

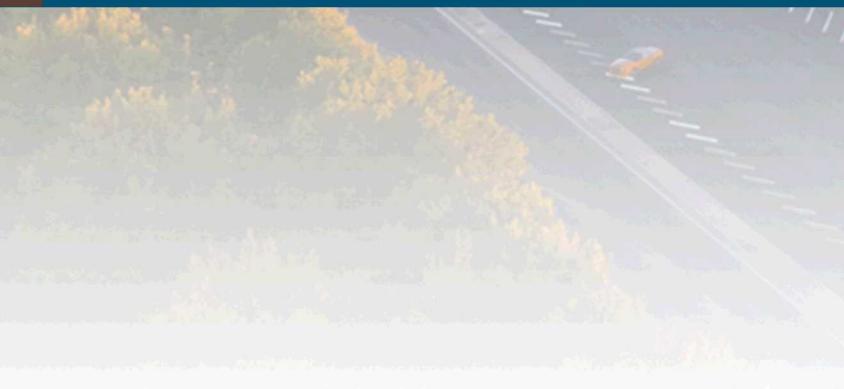


Parameter sweep of SM clock frequency and memory frequency

- SM frequency: 500MHz, **1312MHz**, 1800Mhz (2.5x → 1.3x speedup)
- HBM frequency: **877MHz**, 1200MHz, 1600Mhz (0.99x → 1.04x speedup)



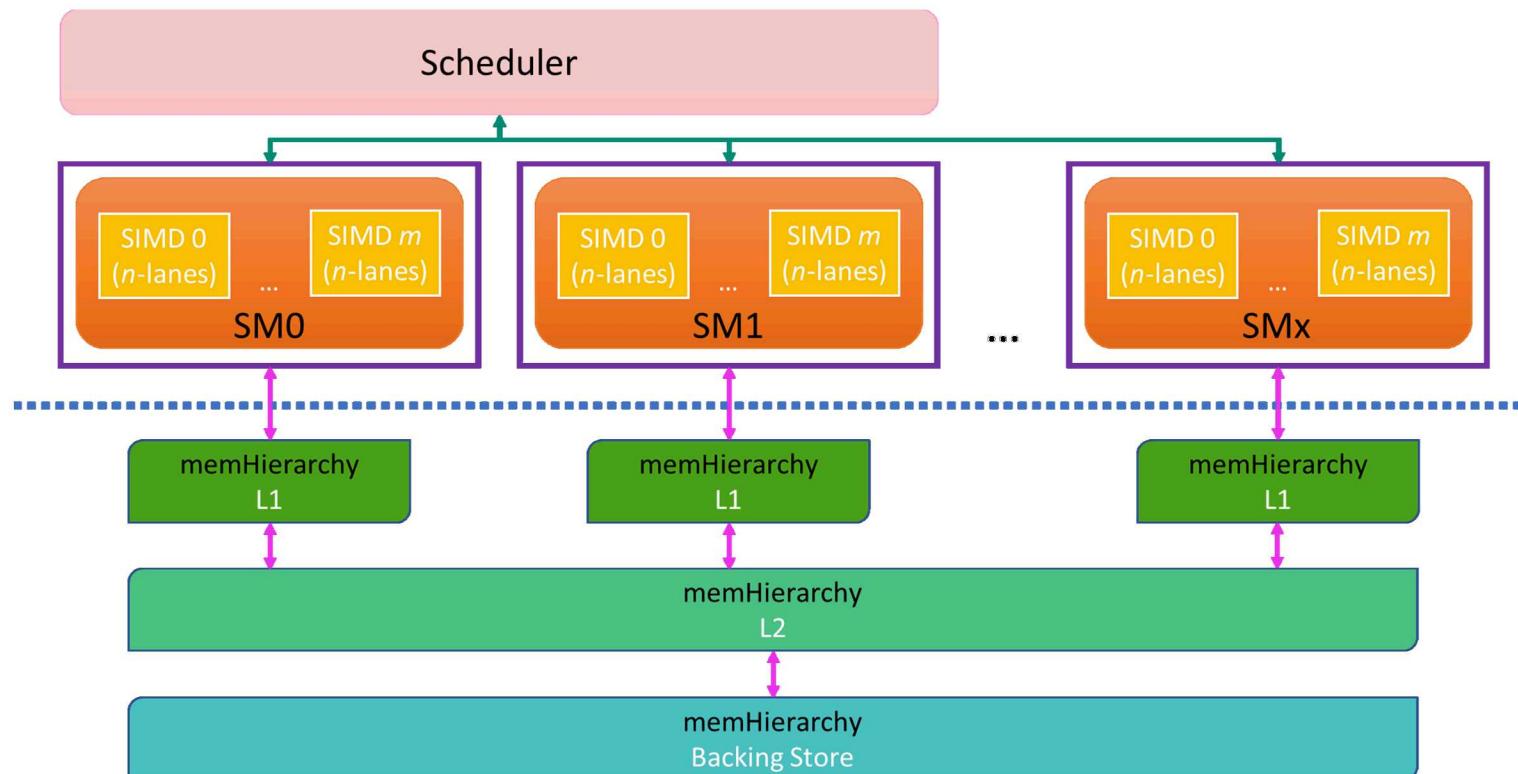
# Future Work and Summaries



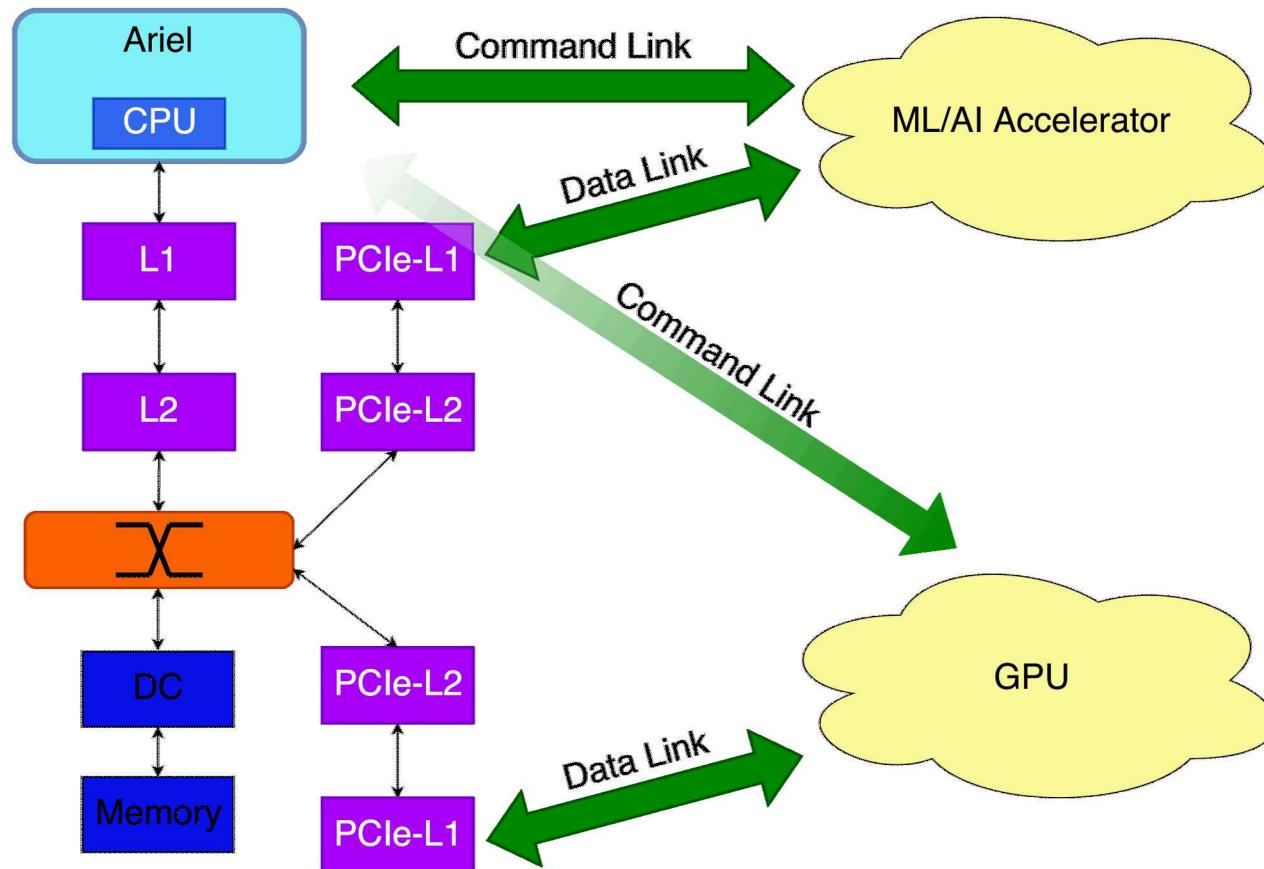
# Future Work

## Split Monolithic GPGPU Component into SM Components

- Speed up simulations
- Wider range of system options
  - Chiplets
  - Split interconnects

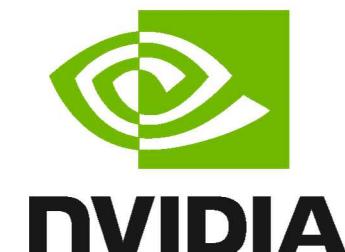
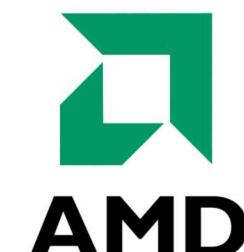
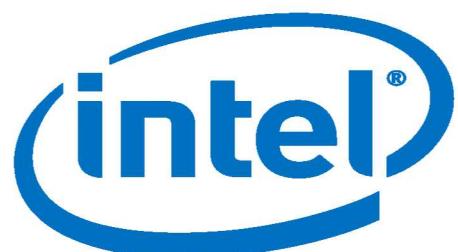


# Position SST For the Future



Proves that SST can be integrated with complex external models

- Exploring new ISA plugins for Ariel
  - ARM (SVE?), other alternatives
  - HDL simulation components
  - Low-level hardware design including possible path to tape-out
- Neural-inspired & quantum
  - Exploration of interfaces and data movement
- Near memory/storage computing
  - Micron
- Easier integration with vendor-supplied accelerator models



# Lessons Learned



The philosophy behind the design of SST has always been modularity

- Any component should be able to communicate with any other component, assuming that the appropriate interface is available, with no other dependencies (mostly)
- The integration of the GPU component has validated this design approach
  - Dropped in an “unknown” component and the memory subsystem...works!

The academic community, particularly in architecture, has little to no exposure to applications of ASC-level code complexity

- Genuinely surprised that so many of the applications had trouble with the parser
- Some instructions not even implemented because they had never encountered them

SST needs an “accelerator” component to hide complexity

- Will probably look like the current GPU model with bus links (PCIe, Gen-Z, CXL) and one or more back channel command/control links



# Remaining Issues

## Support for UVM

Statistics are still resident in the GPU component and separate from SST

PTX parser has trouble with some instructions when compiling for sm\_70

PTX parser is unable to locate a post dominator instruction in some instances

Texture cache is not implemented in functional and timing models

- Presents a problem for some Kokkos-based applications

Some of the Thrust library is not implemented or is implemented incorrectly leading to segfaults

Need write-through support in memHierarchy

Upstream GPGPUSim changes to mainline

## Next Demo...

Full day tutorial (with hands-on) at IISWC 2019 [<http://www.iiswc.org/iiswc2019>]

- November 03





# Primary SST Components: memHierarchy

Collection of interoperable memory system elements

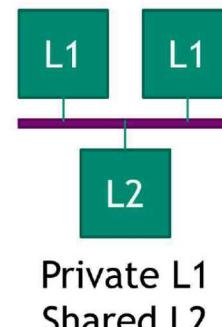
- Caches and directories
- Memory models (DDR, HBM, HMC, NVRAM, etc.)
- Memory controllers & network interfaces for memory (MemNICs)

Inter- and intra-socket coherence

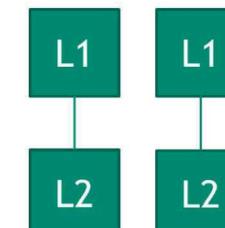
Correlated with modern memory hierarchies

- HBM2/3 Evaluation on Many-core CPU
- Messier: A Detailed NVM-Based DIMM Model for the SST Simulation Framework

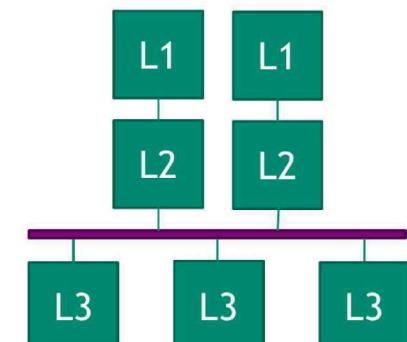
Will allow us to model any sort of memory arrangement that we can imagine!



Private L1  
Shared L2



Private L1  
Private L2  
*(would require directory for coherence)*



Private L1  
Private L2  
Distributed shared L3



## PTX

```

$Lt_25_13570:
    ld.global.s32    %r9, [%rd5+0];
    add.s32          %r10, %r9, %r8;
    ld.global.s32    %r11, [%rd5+1024];
    add.s32          %r8, %r11, %r10;
    add.u32          %r5, %r7, %r5;
    add.u64          %rd5, %rd5, %rd6;
    ld.param.u32    %r6, [size];
    setp.lt.u32    %p2, %r5, %r6;
    @%p2 bra        $Lt_25_13570;

...
    mov.u32 %r12, 127;
    setp.gt.u32    %p3, %r3, %r12;
    @%p3 bra        $Lt_25_14082;
    ld.shared.s32   %r13, [%rd10+512];
    add.s32          %r8, %r13, %r8;
    st.shared.s32   [%rd10+0], %r8;

$Lt_25_14082:
    bar.sync 0;

```

## SASS (PTXPlus)

```

l0x00000060:
    add.half.u32    $r7, $r4, 0x00000400;
    ld.global.u32   $r8, [$r4];
    ld.global.u32   $r7, [$r7];
    add.half.u32    $r0, $r5, $r0;
    add.half.u32    $r6, $r8, $r6;
    set.gt.u32.u32  $p0/$o127, s[0x0020], $r0;
    add.half.u32    $r6, $r7, $r6;
    add.half.u32    $r4, $r4, $r3;
    @$p0.ne bra    l0x00000060;

...
    set.gt.u32.u32  $p0/$o127, $r2, const [0x0000];
    @$p0.equ add.u32 $ofs2, $ofs1, 0x00000230;
    @$p0.equ add.u32 $r6, s[$ofs2+0x0000], $r6;
    @$p0.equ mov.u32  s[$ofs1+0x0030], $r6;
    bar.sync 0x00000000;

```

# Why SST?

Problem: Simulation is slow

- Tradeoff between accuracy and time to simulate
- Many simulators are serial, unable to simulate very large systems

Problem: Lack of simulator flexibility

- Tightly-coupled simulations: Difficult to modify
- Difficult to simulate at different levels of accuracy

The Structural Simulation Toolkit:  
A parallel, discrete-event simulation framework  
for scalability and flexibility

# What is SST?

<b>Goals</b>	<b>Status</b>
<ul style="list-style-type: none"> <li>• Become the standard architectural simulation framework for HPC</li> <li>• Be able to evaluate future systems on DOE/DOD workloads</li> <li>• Use supercomputers to design supercomputers</li> </ul>	<ul style="list-style-type: none"> <li>• Parallel Core, basic components</li> <li>• Current Release (7.1) <ul style="list-style-type: none"> <li>• Improved components</li> <li>• Modular core/elements</li> <li>• More Internal documentation</li> </ul> </li> </ul>
<b>Technical Approach</b>	<b>Consortium</b>
<ul style="list-style-type: none"> <li>• Parallel <ul style="list-style-type: none"> <li>• Parallel Discrete Event core with <u>conservative optimization over MPI/Threads</u></li> </ul> </li> <li>• Multiscale <ul style="list-style-type: none"> <li>• Detailed and simple models for processor, network, &amp; memory</li> </ul> </li> <li>• Interoperability <ul style="list-style-type: none"> <li>• DRAMSim, ,memory models</li> <li>• routers, NICs, schedulers</li> </ul> </li> <li>• Open <ul style="list-style-type: none"> <li>• Open Core, non-viral, modular</li> </ul> </li> </ul>	<p>The image displays a grid of logos representing various SST Consortium members. The logos are arranged in three rows. The first row contains the University of Delaware logo (blue 'U' shape) and the University of Maryland logo (Maryland state map). The second row contains the NVIDIA logo (green stylized eye), the gem5 logo (blue '5' with a grey 'm'), the AWE logo (blue atom symbol), the hp logo (blue 'hp' monogram), the ND logo (blue 'ND' monogram), the Mellanox Technologies logo (purple triangle), the CRAY logo (blue 'CRAY' monogram), the Micron logo (blue 'Micron' monogram), and the Oak Ridge National Laboratory logo (green leaf). The third row contains the BOSTON UNIVERSITY logo (red rectangle with 'BOSTON UNIVERSITY' text), the Intel logo (blue 'intel' monogram), and the Wisconsin logo (University of Wisconsin-Madison crest).</p>

# Key Capabilities

## Parallel

- Built from the ground up to be scalable
- Demonstrated scaling to 512+ host processors
- Conservative, Distance-based Optimization
- MPI + Threads

## Flexible

- Enables “mix and match” of simulation components
- Multiscale tradeoff between accuracy and simulation time
  - e.g., cycle-accurate network with trace-driven endpoints
- Open API
  - Easily extensible with new models
  - Modular framework
  - Open-source core

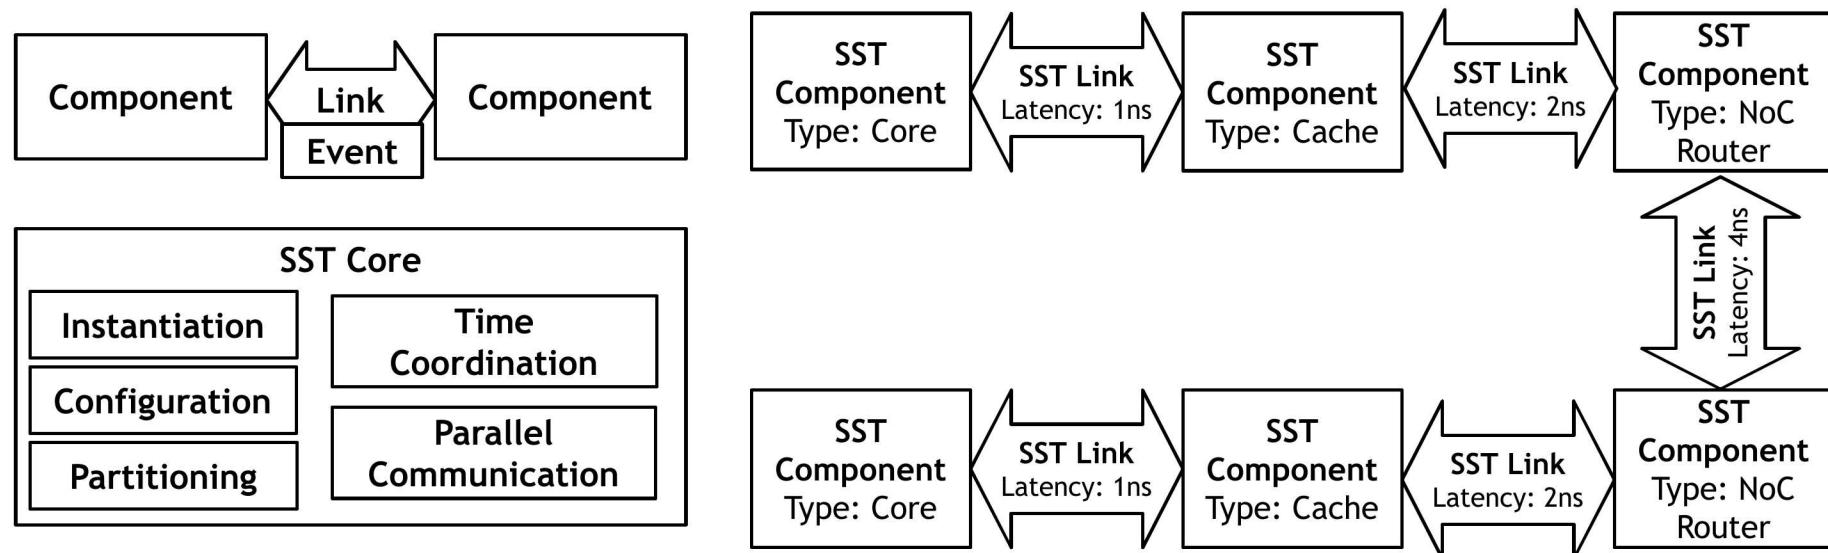
# SST's discrete-event algorithm

Simulations are comprised of **components** connected by **links**

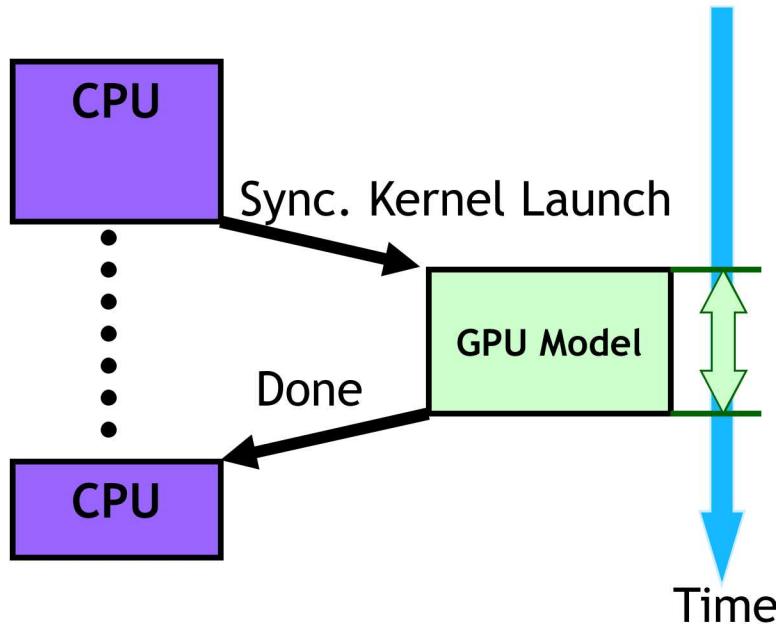
**Components** interact by sending events over **links**

Each **link** has a minimum latency

**Components** can load **subComponents** and **modules** for additional functionality



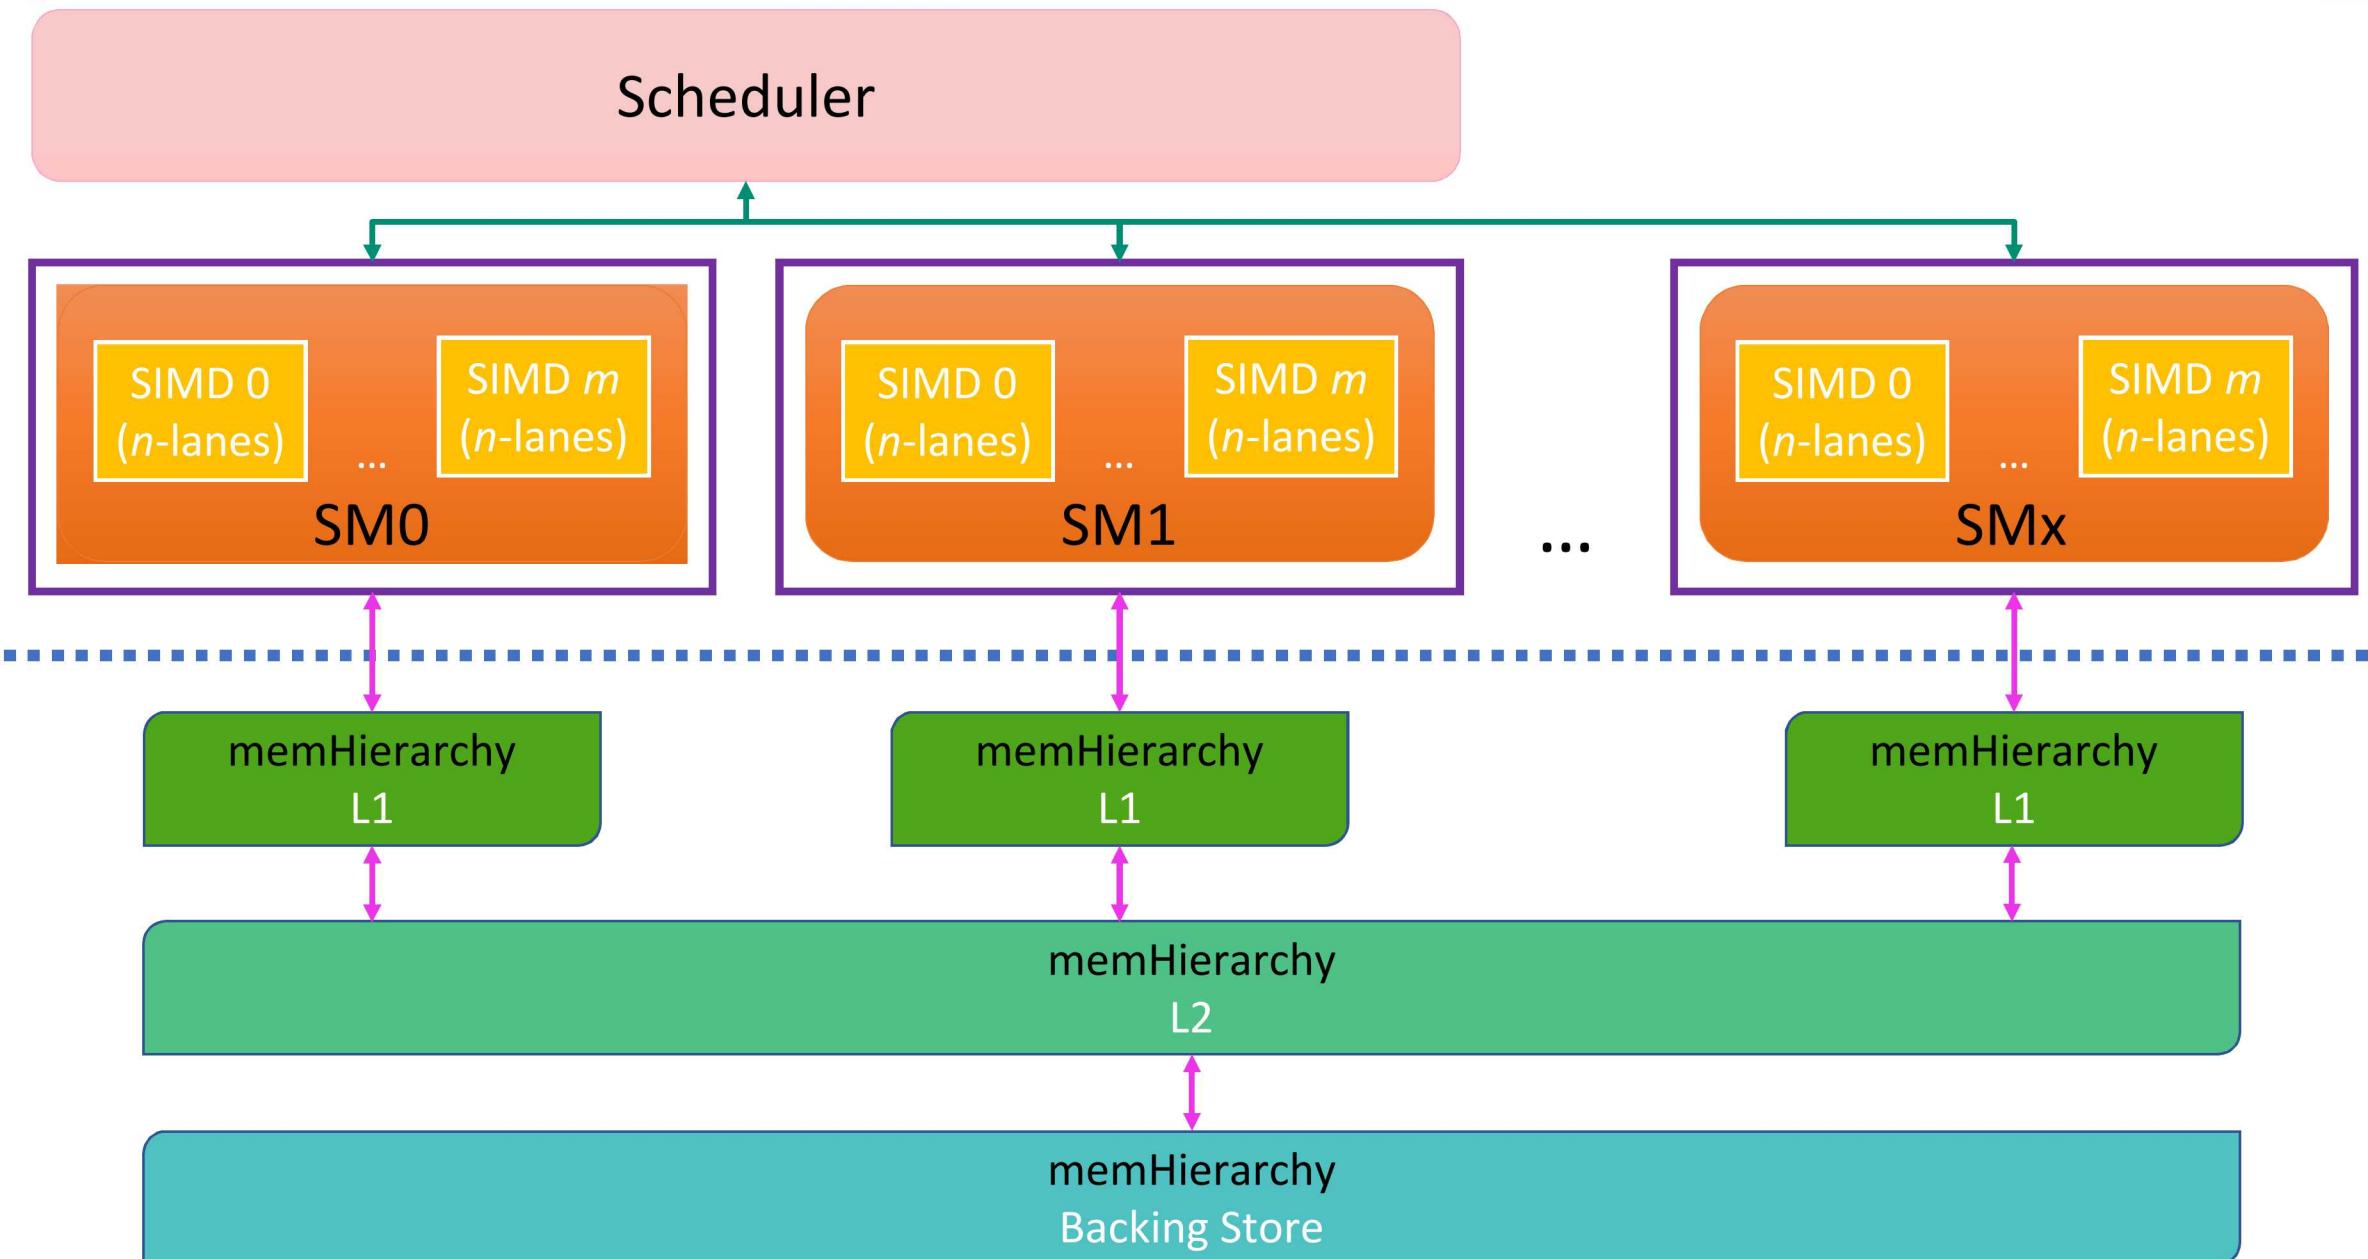
Currently only supports synchronous kernel launches



# Kokkos Kernels Unit Test Results

1	abs_double	34	sparse_gauss_seidel_double_int_int_TestExecSpace	67	batched_scalar_serial_trsm_l_l_t_u_double_double
2	abs_mv_double	35	sparse_gauss_seidel_double_int_size_t_TestExecSpace	68	batched_scalar_serial_trsm_l_l_t_n_double_double
3	asum_double	36	sparse_block_gauss_seidel_double_int_int_TestExecSpace	69	batched_scalar_serial_trsm_l_u_t_u_double_double
4	axpby_double	37	sparse_block_gauss_seidel_double_int_size_t_TestExecSpace	70	batched_scalar_serial_trsm_l_u_t_n_double_double
5	axpby_mv_double	38	sparse_crsmatrix_double_int_int_TestExecSpace	71	batched_scalar_serial_gemv_nt_double_double
6	axpy_double	39	sparse_crsmatrix_double_int_size_t_TestExecSpace	72	batched_scalar_serial_gemv_t_double_double
7	axpy_mv_double	40	sparse_blkcrsmatrix_double_int_int_TestExecSpace	73	batched_scalar_serial_trsv_l_nt_u_double_double
8	dot_double	41	sparse_blkcrsmatrix_double_int_size_t_TestExecSpace	74	batched_scalar_serial_trsv_l_nt_n_double_double
9	dot_mv_double	42	sparse_replaceSumIntoLonger_double_int_int_TestExecSpace	75	batched_scalar_serial_trsv_u_nt_u_double_double
10	mult_double	43	sparse_replaceSumIntoLonger_double_int_size_t_TestExecSpace	76	batched_scalar_serial_trsv_u_nt_n_double_double
11	mult_mv_double	44	sparse_replaceSumInto_double_int_int_TestExecSpace	77	batched_scalar_team_set_double_double
12	nrm1_double	45	sparse_replaceSumInto_double_int_size_t_TestExecSpace	78	batched_scalar_team_scale_double_double
13	nrm1_mv_double	46	graph_graph_color_double_int_int_TestExecSpace	79	batched_scalar_team_gemm_nt_nt_double_double
14	nrm2_double	47	graph_graph_color_double_int_size_t_TestExecSpace	80	batched_scalar_team_gemm_t_nt_double_double
15	nrm2_mv_double	48	graph_graph_color_deterministic_double_int_int_TestExecSpace	81	batched_scalar_team_gemm_nt_t_double_double
16	nrm2_squared_double	49	graph_graph_color_deterministic_double_int_size_t_TestExecSpace	82	batched_scalar_team_gemm_t_t_double_double
17	nrm2_squared_mv_double	50	graph_graph_color_d2_double_int_int_TestExecSpace	83	batched_scalar_team_trsm_l_l_nt_u_double_double
18	nrminf_double	51	graph_graph_color_d2_double_int_size_t_TestExecSpace	84	batched_scalar_team_trsm_l_l_nt_n_double_double
19	nrminf_mv_double	52	common_ArithTraits	85	batched_scalar_team_trsm_l_u_nt_u_double_double
20	reciprocal_double	53	common_set_bit_count	86	batched_scalar_team_trsm_l_u_nt_n_double_double
21	reciprocal_mv_double	54	common_ffs	87	batched_scalar_team_trsm_r_u_nt_u_double_double
22	scal_double	55	batched_scalar_serial_set_double_double	88	batched_scalar_team_trsm_r_u_nt_n_double_double
23	scal_mv_double	56	batched_scalar_serial_scale_double_double	89	batched_scalar_team_trsm_l_l_t_u_double_double
24	sum_double	57	batched_scalar_serial_gemm_nt_nt_double_double	90	batched_scalar_team_trsm_l_l_t_n_double_double
25	sum_mv_double	58	batched_scalar_serial_gemm_t_nt_double_double	91	batched_scalar_team_trsm_l_u_t_u_double_double
26	update_double	59	batched_scalar_serial_gemm_nt_t_double_double	92	batched_scalar_team_trsm_l_u_t_n_double_double
27	update_mv_double	60	batched_scalar_serial_gemm_t_t_double_double	93	batched_scalar_team_gemv_nt_double_double
28	gemv_double	61	batched_scalar_serial_trsm_l_l_nt_u_double_double	94	batched_scalar_team_gemv_t_double_double
29	gemm_double	62	batched_scalar_serial_trsm_l_l_nt_n_double_double	95	batched_scalar_serial_lu_double
30	sparse_spgemm_double_int_int_TestExecSpace	63	batched_scalar_serial_trsm_l_u_nt_u_double_double	96	batched_scalar_serial_inverselu_double
31	sparse_spgemm_double_int_size_t_TestExecSpace	64	batched_scalar_serial_trsm_l_u_nt_n_double_double	97	batched_scalar_serial_solvlu_double
32	sparse_spadd_double_int_int_TestExecSpace	65	batched_scalar_serial_trsm_r_u_nt_u_double_double	98	batched_scalar_team_lu_double
33	sparse_spadd_double_int_size_t_TestExecSpace	66	batched_scalar_serial_trsm_r_u_nt_n_double_double	99	batched_scalar_team_inverselu_double

# Split Monolithic GPGPU Component into SM Components

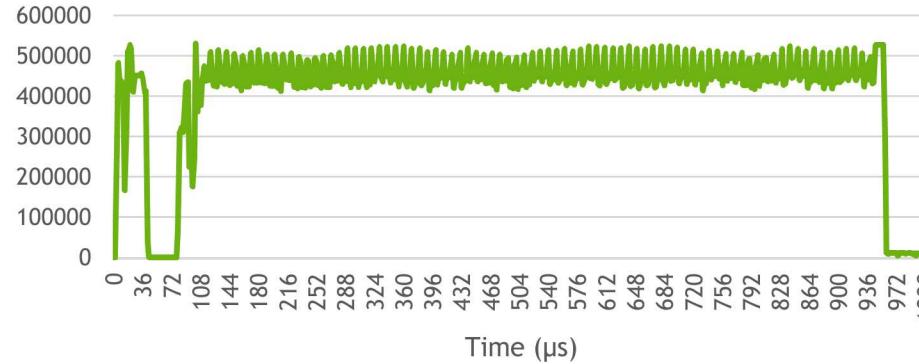


# Performance Analysis – Lulesh

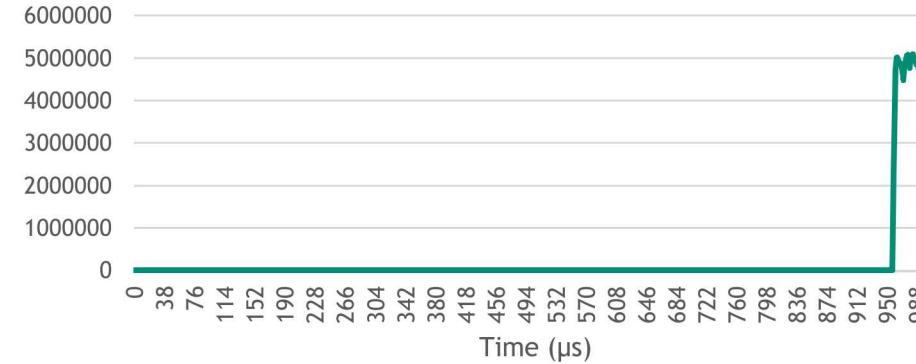
SST is capable of providing periodic statistic dumps for all of the currently loaded components

- Valuable for fine-grain performance analysis of applications

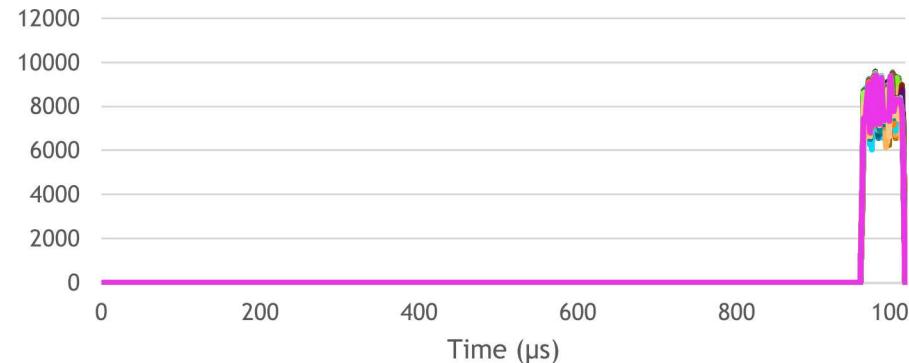
Host Cycles



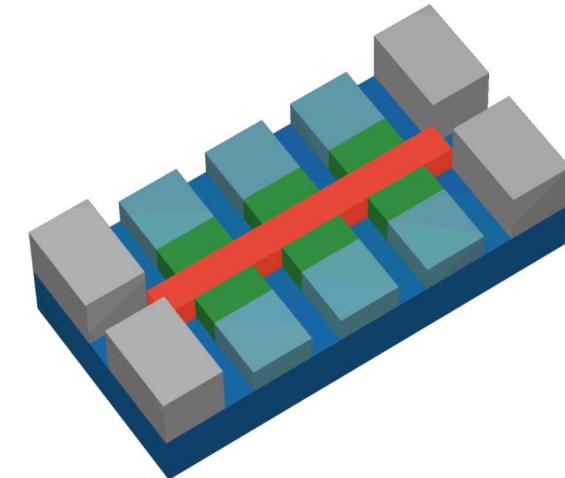
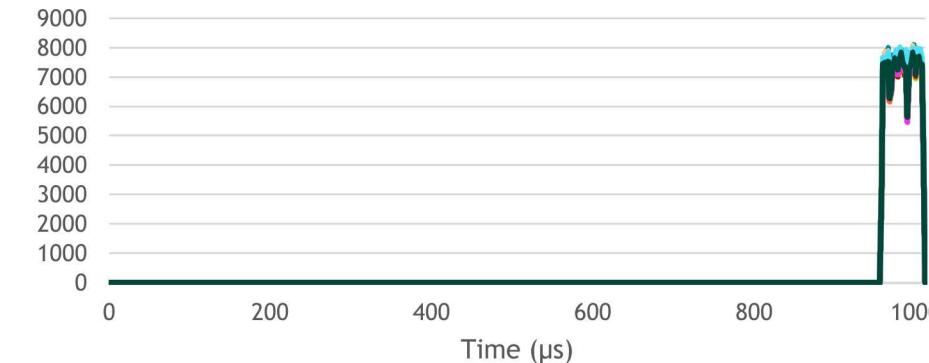
Crossbar Packets Processed



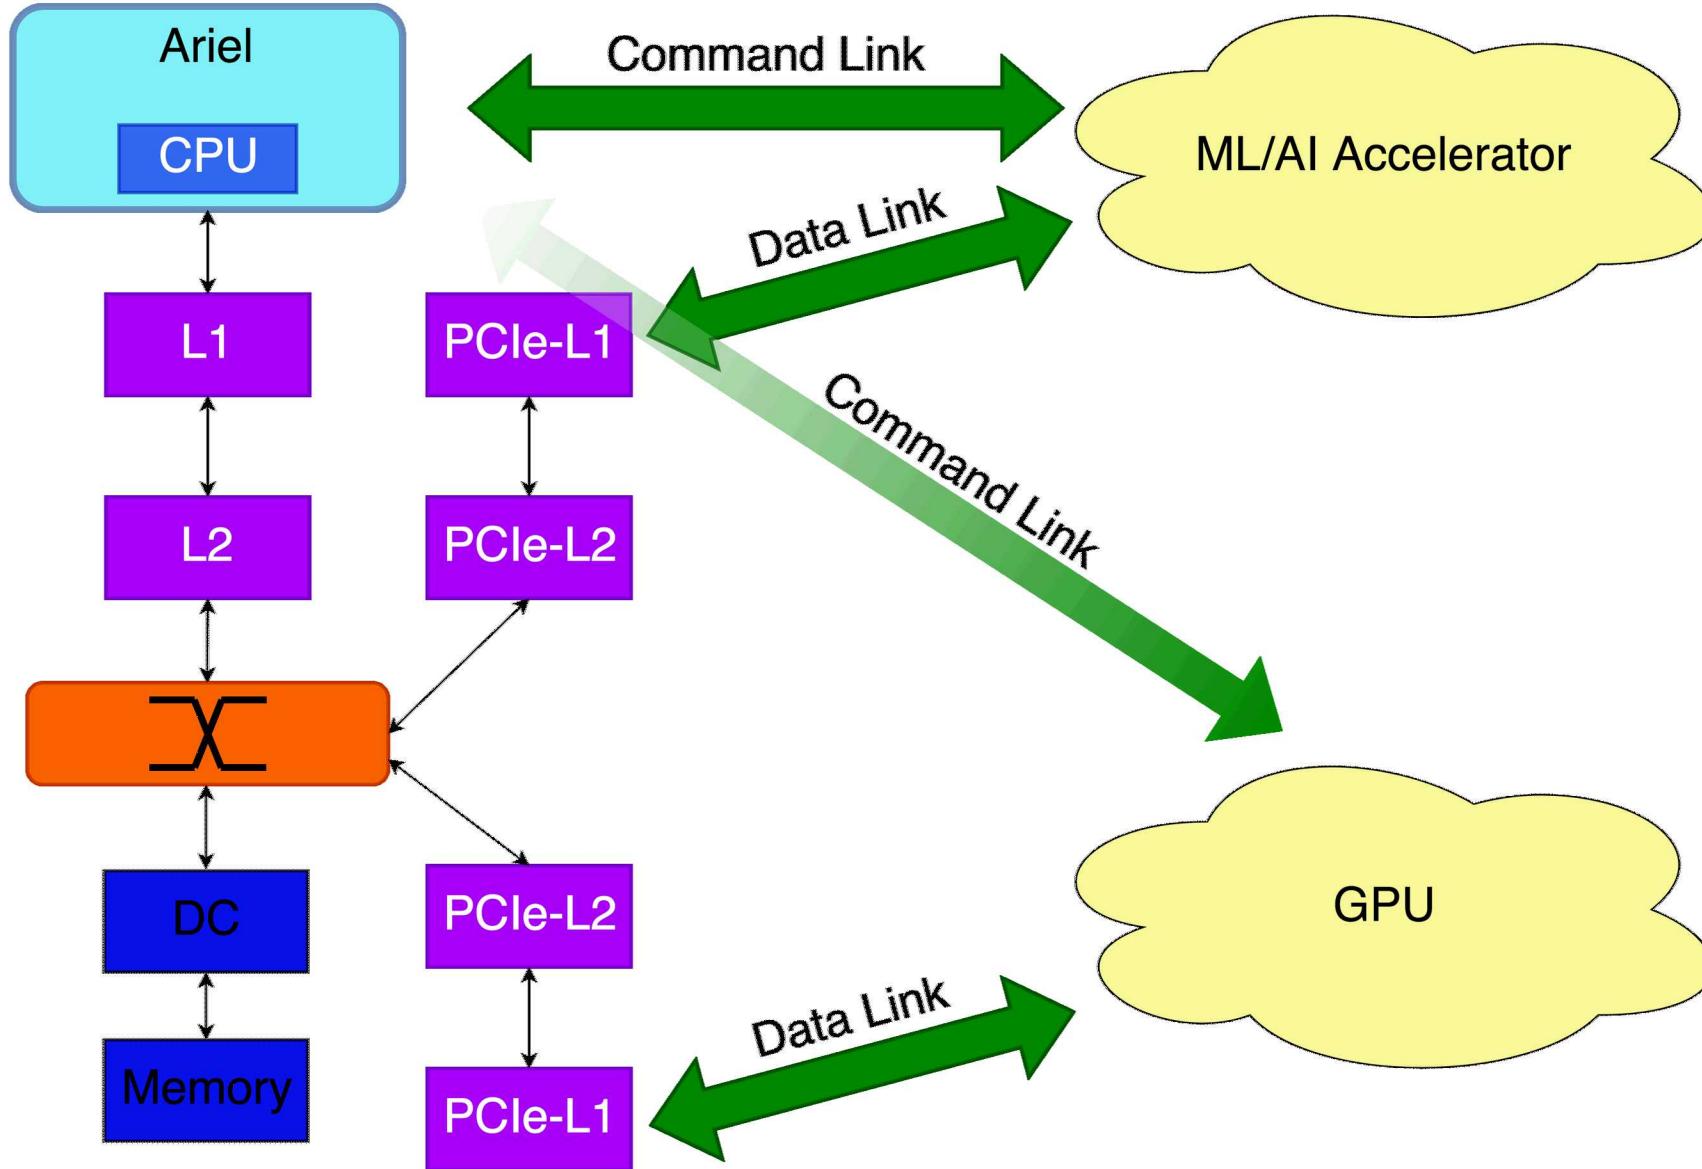
L1\$ Hits



L2\$ Hits



# GPGPUSim Integration



CLICK TO EDIT MASTER TITLE STYLE



Slide Left Blank